

1 **TITLE:** A phylogenetic transform enhances analysis of compositional microbiota data
2

3 **AUTHORS:** Justin D Silverman^{*†‡}, Alex D Washburne^{¶#}, Sayan Mukherjee^{*||}, and
4 Lawrence A David^{*‡§}
5

6 **AUTHOR AFFILIATIONS:**

7 * Program in Computational Biology and Bioinformatics, Duke University, Durham, NC
8 27708

9 † Medical Scientist Training Program, Duke University, Durham, NC 27708

10 ‡ Center for Genomic and Computational Biology, Duke University, Durham, NC 27708

11 § Department of Molecular Genetics and Microbiology, Duke University, Durham, NC
12 27708

13 ¶ Nicholas School of the Environment, Duke University, Durham, NC 27708

14 # Cooperative Institute for Research in Environmental Sciences (CIRES), University of
15 Colorado, Boulder, CO 80309

16 || Departments of Statistical Science, Mathematics, and Computer Science, Duke
17 University, Durham, NC 27708

18

19 **CORRESPONDING AUTHOR:**

20 Lawrence A David

21 CIEMAS, Room 2171

22 101 Science Drive, Box 3382

23 Durham, NC 27708

24 919-668-5388

25 lawrence.david@duke.edu

26

27

28 **ABSTRACT:** Surveys of microbial communities (microbiota), typically measured as
29 relative abundance of species, have illustrated the importance of these communities in
30 human health and disease. Yet, statistical artifacts commonly plague the analysis of
31 relative abundance data. Here, we introduce the PhILR transform, which incorporates
32 microbial evolutionary models with the isometric log-ratio transform to allow off-the-shelf
33 statistical tools to be safely applied to microbiota surveys. We demonstrate that
34 analyses of community-level structure can be applied to PhILR transformed data with
35 performance on benchmarks rivaling or surpassing standard tools. Additionally, By
36 decomposing distance in the PhILR transformed space, we identified neighboring
37 clades that may have adapted to distinct human body sites. Decomposing variance
38 revealed that covariation of bacterial clades within human body sites increases with
39 phylogenetic relatedness. Together, these findings illustrate how the PhILR transform
40 combines statistical and phylogenetic models to overcome compositional data
41 challenges and enable evolutionary insights relevant to microbial communities.

42

43 **IMPACT STATEMENT:** The PhILR transform uses an evolutionary model to overcome
44 statistical challenges associated with microbiota surveys.

45

46 **INTRODUCTION**

47

48 Microbiota research today embodies the data-rich nature of modern biology.

49 Advances in high-throughput DNA sequencing allow for rapid and affordable surveys of

50 thousands of bacterial taxa across hundreds of samples (1). The exploding availability

51 of sequencing data has poised microbiota research to advance our understanding of

52 fields as diverse as ecology, evolution, medicine, and agriculture (2). Considerable

53 effort now focuses on interrogating microbiota datasets to identify relationships between

54 bacterial taxa, as well as between microbes and their environment.

55 Increasingly, it is appreciated that the relative nature of microbial abundance data

56 in microbiota studies can lead to spurious statistical analyses (3-9). With next

57 generation sequencing, the number of reads per sample can vary independently of

58 microbial load (6, 9). In order to make measurements comparable across samples, most

59 studies therefore analyze the relative abundance of bacterial taxa. Analyses are thus

60 not carried out on absolute abundances of community members (Figure 1A), but rather

61 on relative data occupying a constrained geometric space and represented in a non-

62 Cartesian coordinate system (Figure 1B). Such relative abundance datasets are often

63 termed compositional. The use of most standard statistical tools (e.g., correlation,

64 regression, or classification) within a compositional space leads to spurious results (10).

65 For example, three-quarters of the significant bacterial interactions inferred by Pearson

66 correlation on a compositional human microbiota dataset were likely false (4), and over

67 two-thirds of differentially abundant taxa inferred by a t-test on a simulated

68 compositional human microbiota dataset were spurious (11). To account for

69 compositional effects in microbial datasets, bioinformatics efforts have re-derived

70 common statistical methods including correlation statistics (4, 12), hypothesis testing
71 (13), and variable selection (14, 15).

72 An alternative approach is to transform compositional microbiota data to a space
73 where existing statistical methods may be applied without introducing spurious
74 conclusions. This approach is attractive because of its efficiency: the vast toolbox of
75 existing statistical models can be applied without re-derivation. Normalization methods,
76 for example, have been proposed to modify count data by assuming reads follow certain
77 statistical distributions (e.g., negative binomial) (16, 17). Alternatively, the field of
78 Compositional Data Analysis (CoDA) has focused on formalizing methods for
79 transforming compositional data into a simpler geometry without having to assume data
80 adhere to a distribution model (18). Previous microbiota analyses have already
81 leveraged CoDA theory and used the centered log-ratio transform to reconstruct
82 microbial association networks and interactions (19, 20) and to analyze differential
83 abundances (21, 22). However, the centered log-ratio transform has a crucial limitation:
84 it yields a coordinate system featuring a singular covariance matrix and is thus
85 unsuitable for many common statistical models (10). This drawback can be sidestepped
86 using another CoDA transform, known as the Isometric Log-Ratio (ILR) transformation
87 (23). The ILR transform can be built from a sequential binary partition of the original
88 variable space (Figure 1C), creating a new coordinate system with an orthonormal basis
89 (Figure 1D and E) (24). However, a known obstacle to using the ILR transform is the
90 choice of partition such that the resulting coordinates are meaningful (10). To date,
91 microbiota studies have chosen ILR coordinates using ad hoc sequential binary

92 partitions of bacterial groups that are not easily interpreted (25, 26). Alternatively,
93 external covariates have been used to pick groups of bacterial taxa to contrast (27).

94 Here, we introduce the bacterial phylogenetic tree as a natural and informative
95 sequential binary partition when applying the ILR transform to microbiota datasets
96 (Figure 1C). Using phylogenies to construct the ILR transform results in an ILR
97 coordinate system capturing evolutionary relationships between neighboring bacterial
98 groups (clades). Analyses of neighboring clades offer the opportunity for biological
99 insight: clade analyses have linked genetic differentiation to ecological adaptation (28),
100 and the relative levels of sister bacterial genera differentiate human cohorts by diet,
101 geography, and culture (29-31). Datasets analyzed by a phylogenetically aware ILR
102 transform could therefore reveal ecological and evolutionary factors shaping host-
103 associated microbial communities.

104 We term our approach the **Phylogenetic ILR** (PhILR) transform. Using published
105 environmental and human-associated 16S rRNA datasets as benchmarks, we found
106 that simple Euclidean distances calculated on PhILR transformed data provided a
107 compositionally robust alternative to distance/dissimilarity measures like Bray-Curtis,
108 Jaccard, and Unifrac. In addition, we observed that the accuracy of supervised
109 classification methods on our benchmark datasets was matched or improved with
110 PhILR transformed data relative to applying the same models on untransformed (raw) or
111 log transformed relative abundance data. Decomposing distances between samples
112 along PhILR coordinates identified bacterial clades that may have differentiated to adapt
113 to distinct body sites. Similar decomposition of variance along PhILR coordinates
114 showed that, in all human body sites studied, the degree to which neighboring bacterial

115 clades covary tends to increase with the phylogenetic relatedness between clades.
116 Together, these findings demonstrate that the PhILR transform can be used to enhance
117 existing microbiota analysis pipelines, as well as enable novel phylogenetic analyses of
118 microbial ecosystems.

119

120 RESULTS

121 Constructing the PhILR transform

122 The PhILR transform has two goals. The first goal is to transform input microbiota
123 data into an unconstrained space with an orthogonal basis while preserving all
124 information contained in the original composition. The second goal is to conduct this
125 transform using phylogenetic information. To achieve these dual goals on a given set of
126 N samples consisting of relative measurements of D taxa (Figure 1B), we transform
127 data into a new space of N samples and $(D - 1)$ coordinates termed ‘balances’ (Figure
128 1C-E) (23, 24). Each balance y_i^* is associated with a single internal node i of a
129 phylogenetic tree with the D taxa as leaves (the asterisk denotes a quantity represented
130 in PhILR space). The balance represents the log-ratio of the geometric mean relative
131 abundance of the two clades of taxa that descend from i (*Methods*). Although individual
132 balances may share overlapping sets of leaves and thus exhibit dependent behavior,
133 the ILR transform rescales and combines leaves to form a coordinate system whose
134 basis vectors are orthonormal and the corresponding coordinates are Cartesian (23,
135 24). The orthogonality of basis vectors allows conventional statistical tools to be used
136 without compositional artifacts. The unit-length of basis vectors makes balances across
137 the tree statistically comparable even when they have differing numbers of descendant

138 tips or exist at different depths in the tree (10). In addition, the unit-length ensures that
139 the variance of PhILR balances has a consistent scale, unlike the variance of log-ratios
140 originally proposed by Aitchison (5) as a measure of association, in which it can be
141 unclear what constitutes a large or small variance (4).

142 While the above description represents the core of the PhILR transform, we have
143 also equipped the PhILR transform with two sets of weights that can: 1) address the
144 multitude of zero and near-zero counts present in microbiota data; and, 2) incorporate
145 phylogenetic branch lengths into the transformed space. Because zero counts cause
146 problems when computing logs or performing division, zeros are often replaced in
147 microbiota analyses with small non-zero counts. However, to avoid excess zero
148 replacement that may itself introduce bias, stringent hard filtering thresholds are often
149 employed (e.g. removing all taxa that are not seen with at least a minimum number of
150 counts in a subset of samples). Still, hard filtering thresholds may remove a substantial
151 fraction of observed OTUs and do not account for the low precision (or high variability)
152 of near-zero counts (7, 32, 33). We therefore developed a ‘taxon weighting’ scheme that
153 acts as a type of soft-thresholding, supplementing zero replacement methods with a
154 generalized form of the ILR transform that allows weights to be attached to individual
155 taxa (34). Weights are chosen with a heuristic designed to down weight the influence of
156 taxa with many zero or near-zero counts (*Methods*).

157 Our second weighting scheme is called branch length weighting. Certain
158 analyses may benefit from incorporating information on evolutionary distances between
159 taxa (35-37). For example, because related bacteria may be more likely to share similar
160 traits (38), it may be desirable to consider communities differing only in the abundance

161 of closely-related microbes to be more similar than communities differing only in the
162 abundance of distantly-related microbes. Because of the one-to-one correspondence
163 between PhILR balances and internal nodes on the phylogenetic tree, evolutionary
164 information can be incorporated into the PhILR transform by scaling balances using the
165 phylogenetic distance between their direct descendants (*Methods*). We note that we
166 employ both branch length weighting and taxon weighting throughout our following
167 analyses except where noted; still, these weights should be considered optional
168 additions to the core PhILR transform.

169

170 **Benchmarking community-level analyses in the PhILR coordinate system**

171 To illustrate how the PhILR transform can be used to perform standard
172 community-level analyses of microbiota datasets, we first examined measures of
173 community dissimilarity. Microbiota analyses commonly compute the dissimilarity or
174 distance between pairs of samples and use these computed pairwise distances as input
175 to a variety of statistical tools. We investigated how Euclidean distances calculated on
176 PhILR transformed data compared to common ecological measures of microbiota
177 distance or dissimilarity (UniFrac, Bray-Curtis, and Jaccard) as well as Euclidean
178 distance applied to raw relative abundance data in standard distance-based analysis.
179 We chose three different microbiota surveys as reference datasets: Costello Skin Sites
180 (CSS), a dataset of 357 samples from 12 human skin sites (39, 40); Human Microbiome
181 Project (HMP), a dataset of 4,743 samples from 18 human body sites (e.g., skin,
182 vaginal, oral, and stool) (41); and, Global Patterns (GP), a dataset of 26 samples from

183 9 human or environmental sites (1) (Supplementary file 1 and Figure 2-figure
184 supplement 1).

185 Distance-based analyses using Euclidean distances computed on PhILR
186 transformed data exhibited performance rivaling common ecological distance or
187 dissimilarity measures. Principal coordinate analyses (PCoA) qualitatively demonstrated
188 separation of body sites using both Euclidean distances on PhILR transformed data
189 (Figure 2A) and with several standard distance measures calculated on raw relative
190 abundance data (Figure 2-figure supplement 2). To quantitatively compare distance
191 measures, we tested how well habitat information explained variability among distance
192 matrices as measured by the R^2 statistic from PERMANOVA (42). By this metric, the
193 Euclidean distance in the PhILR coordinate system significantly outperformed the five
194 competing distance metrics in all but 1 case (in comparison to Weighted UniFrac when
195 applied to the HMP dataset; Figure 2B).

196 Next, we tested the performance of predictive statistical models in the PhILR
197 coordinate system. We examined four standard supervised classification techniques:
198 logistic regression (LR), support vector machines (SVM), k-nearest neighbors (kNN),
199 and random forests (RF) (40). We applied these methods to the same three reference
200 datasets used in our comparison of distance metrics. As a baseline, the machine
201 learning methods were applied to raw relative abundance datasets and raw relative
202 abundance data that had been log-transformed.

203 The PhILR transform significantly improved supervised classification accuracy in
204 7 of the 12 benchmark tasks compared to raw relative abundances (Figure 2C).
205 Accuracy improved by more than 90% in two benchmarks (SVM on HMP and GP),

206 relative to results on the raw data. Log transformation of the data also improved
207 classifier accuracy significantly on 6 of the 12 benchmarks but also significantly
208 underperformed on 1 benchmark compared to raw relative abundances. In addition, the
209 PhILR transform significantly improved classification accuracy in 5 of the 12
210 benchmarks relative to the log transform. Overall, the PhILR transform often
211 outperformed the raw and log transformed relative abundances with respect to
212 classification accuracy and was never significantly worse.

213

214 **Identifying neighboring clades that differ by body site preference**

215 While our benchmarking experiments demonstrated how PhILR transformed data
216 performed in community-level analyses, we also wanted to explore potential biological
217 insights afforded by the PhILR coordinate system. We therefore investigated how
218 distances decomposed along PhILR balances using a sparse logistic regression model
219 to examine which balances distinguished human body site microbiota in the HMP
220 dataset. Such balances could be used to identify neighboring bacterial clades whose
221 relative abundances capture community-level differences between body site microbiota.
222 Microbial genetic differentiation may be associated with specialization to new resources
223 or lifestyle preferences (28), meaning that distinguishing balances near the tips of the
224 bacterial tree may correspond to clades adapting to human body site environments.

225 We identified dozens of highly discriminatory balances, which were spread
226 across the bacterial phylogeny (Figure 3A and Figure 3-figure supplement 1). Some
227 discriminatory balances were found deep in the tree. Abundances of the Firmicutes,
228 Bacteroidetes, and Proteobacteria relative to the Actinobacteria, Fusobacteria, and

229 members of other phyla, separated skin body sites from oral and stool sites (Figure 3B).
230 Levels of the genus *Bacteroides* relative to the genus *Prevotella* differentiated stool
231 microbiota from other communities on the body (Figure 3C). Notably, values of select
232 balances below the genus level also varied by body site. Relative levels of sister
233 *Corynebacterium* species separated human skin sites from gingival sites (Figure 3D).
234 Species-level balances even differentiated sites in nearby habitats; levels of sister
235 *Streptococcus* species or sister *Actinomyces* species vary depending on specific oral
236 sites (Figure 3E and F). These results show that the decomposition of distances
237 between groups of samples along PhILR balances can be used to highlight ancestral
238 balances that distinguish body site microbiota, as well as to identify more recent
239 balances that may separate species that have adapted to inhabit different body sites.

240

241 **Balance variance and microbiota assembly**

242 As a natural extension of our analysis of how distance decomposes along PhILR
243 balances, we next investigated how balance variance decomposed in the PhILR
244 coordinate system. Balance variance is a measure of association between neighboring
245 bacterial clades. When the variance of a balance between two clades approaches zero,
246 the mean abundance of taxa in each of the two clades will be linearly related and thus
247 covary across microbial habitats (43). By contrast, when a balance exhibits high
248 variance, related bacterial clades exhibit unlinked or exclusionary patterns across
249 samples. Unlike standard measures of association (e.g., Pearson correlation) balance
250 variance is robust to compositional effects (10).

251 Our preliminary investigation demonstrated a striking pattern in which balance
252 variance decreased for balances closer to the tips of the phylogeny and increased for
253 balances nearer to the root. To determine if this observed pattern was not the result of
254 technical artifact, we took the following three steps. First, we omitted branch length
255 weights from the transform as we anticipated that branch lengths may vary non-
256 randomly as a function of depth in the phylogeny. Second, we anticipated that balances
257 near the tips of the phylogeny would be associated with fewer read counts and thus
258 would be more biased by our chosen heuristics for taxon weighting and zero
259 replacement. We therefore omitted taxon weighting, employed more stringent filtering
260 thresholds, and conditioned our calculation of balance variance on non-zero counts
261 rather than using zero-replacement (*Methods*). Third, we combined regression and a
262 permutation scheme to test the null hypothesis that the degree to which neighboring
263 clades covary is independent of the phylogenetic distance between them (*Methods*). By
264 permuting tip labels on the tree, our test generates a restricted subset of random
265 sequential binary partitions that still maintains the count variability (and potential biases
266 due to our zero handling methods) of the observed data.

267 With our modified PhILR analysis in place, we observed significantly decreasing
268 balance variances near the tips of the phylogenetic tree for all body sites in the HMP
269 dataset ($p < 0.01$, permutation test with FDR correction; Figure 4A-F and Figure 4-figure
270 supplements 1-2). Low variance balances predominated near the leaves of the tree.
271 Examples of such balances involved *B. fragilis* species in stool (Figure 4H), *Rothia*
272 *mucilaginosa* species in the buccal mucosa (Figure 4J), and *Lactobacillus* species in the
273 mid-vagina (Figure 4L). By contrast, higher variance balances tended to be more basal

274 on the tree. Three examples of high variance balances corresponded with clades at the
275 order (Figure 4G), family (Figure 4I), and genus (Figure 4K) levels. We also observed
276 that the relationship between balance variance and phylogenetic depth varied at
277 different taxonomic scales. LOESS regression revealed that trends between variance
278 and phylogenetic depth were stronger above the species level than below it (*Methods*;
279 Figure 4D-F and Figure 4-figure supplement 2). Overall, the observed pattern of
280 decreasing balance variance near the tips of the phylogenetic tree suggested that
281 closely related bacteria tend to covary in human body sites.

282

283 **DISCUSSION**

284 There exists a symbiosis between our understanding of bacterial evolution and
285 the ecology of host-associated microbial communities (44). Microbiota studies have
286 shown that mammals and bacteria cospeciated over millions of years (45, 46), and
287 human gut microbes have revealed the forces driving horizontal gene transfer between
288 bacteria (47). Evolutionary tools and theory have been used to explain how cooperation
289 benefits members of gut microbial communities (48), and raise concerns that rising
290 rates of chronic disease are linked to microbiota disruption (49). Here, we aimed to
291 continue building links between microbiota evolution and ecology by designing a data
292 transform that uses phylogenetic models to overcome the challenges associated with
293 compositional data while enabling novel evolutionary analyses.

294 We found that the resulting PhILR coordinate system, at least with respect to the
295 performance metrics chosen, led to significantly improved performance for a variety of
296 community-level analyses now used in microbiota analysis. While these results add

297 credence to our proposed approach, we underscore that we do not find it essential that
298 PhILR demonstrates superior benchmark performance to motivate its use in microbiota
299 analysis. We believe that the need for compositionally robust tools has already been
300 well established (3-9) and intended these benchmarks to showcase the flexibility and
301 utility of working with PhILR transformed data. We also note that for some analyses, a
302 phylogeny-based ILR transform will not outperform an ILR transform built from another
303 sequential binary partition. In fact, in the absence of branch length weights, any random
304 ILR partition would yield equivalent results on our benchmark tasks. Instead, what
305 distinguishes the PhILR transform from other ILR transforms is the interpretability of the
306 transformed coordinates. Balances in PhILR space correspond to speciation events,
307 which can be a source for biological insight.

308 For example, performing regression on PhILR transformed data, enabled us to
309 decompose the distance between bacterial communities onto individual locations on the
310 phylogeny, highlighting balances near the tips of the tree that distinguished human body
311 sites. These balances may reflect functional specialization, as ecological partitioning
312 among recently differentiated bacterial clades could be caused by genetic adaptation to
313 new environments or lifestyles (28). Indeed, among oral body sites, we observed
314 consistent site specificity of neighboring bacterial clades within the genera *Actinomyces*
315 (Figure 3F) and *Streptococcus* (Figure 3E). Species within the *Actinomyces* genera
316 have been previously observed to partition by oral sites (50, 51). Even more
317 heterogeneity has been observed within the *Streptococcus* genus, where species have
318 been identified that distinguish teeth, plaque, mucosal, tongue, saliva, and other oral
319 sites (50, 51). This partitioning likely reflects variation in the anatomy and resource

320 availability across regions of the mouth (50), as well as the kinds of surfaces bacterial
321 strains can adhere to (51).

322 We also observed evidence for potential within-genus adaptation to body sites
323 that has not been previously reported. Within the genus *Corynebacterium*, we found
324 ratios of taxa varied among oral plaques and select skin sites (Figure 3D). Although the
325 genus is now appreciated as favoring moist skin environments, the roles played by
326 individual *Corynebacteria* within skin microbiota remain incompletely understood (52).
327 Precisely linking individual *Corynebacterium* species or strains to body sites is beyond
328 the scope of this study due to the limited taxonomic resolution of 16S rRNA datasets
329 (53, 54). Nevertheless, we believe the PhILR coordinate system may be used in the
330 future to identify groups of related bacterial taxa that have undergone recent functional
331 adaptation.

332 Another example of how the PhILR transform can provide biological insights
333 arose in our analysis of how human microbiota variance decomposes along individual
334 balances. We observed that balances between more phylogenetically related clades
335 were significantly more likely to covary than expected by chance. This pattern could
336 reflect evolutionary and ecological forces structuring microbial communities in the
337 human body. Related bacterial taxa have been hypothesized to have similar lifestyle
338 characteristics (38, 55), and may thus covary in human body sites that favor their
339 shared traits (56, 57). An alternative explanation for the balance variation patterns we
340 observed is that sequencing errors and read clustering artifacts are likely to produce
341 OTUs with similar reference sequences and distributions across samples. While we
342 cannot conclusively rule out this alternative hypothesis, we note that it would not explain

343 why signal for taxa co-variation is strongest for balances at higher taxonomic levels and
344 appears to plateau for balances near or below the species level. A biological
345 explanation for the plateauing signal would be that lifestyle characteristics enabling
346 bacteria to persist in human body sites are conserved among strains roughly
347 corresponding to the same species. Follow-up studies are needed to more conclusively
348 understand how balance variance patterns across the phylogenies can be interpreted
349 from an evolutionary standpoint.

350 Though the methods presented here provide a coherent geometric framework for
351 performing microbiota analysis free from compositional artifacts, future refinements are
352 possible. Specifically, we highlight issues relating to our choice of weights, the handling
353 of zero values, and information loss during count normalization. Both the taxa weights
354 and the branch lengths weights we introduce here may be viewed as preliminary
355 heuristics; future work will likely yield additional weighting schemes, as well as
356 knowledge for when a given weighting scheme should be matched to an analysis task.
357 In the case of supervised machine learning, weighting selection could be optimized as
358 part of the training process. Additionally, if it is important that the transformed data has
359 meaningful numerical coordinates, such that one desires to interpret the exact
360 numerical value of a given balance in a sample, we suggest that neither branch length
361 weights nor taxa weights be used as these weights can complicate this type of
362 interpretation. Concerning our handling of zero values, this model design choice
363 confronts an outstanding challenge for microbiota and compositional data analysis (9,
364 58). Part of this challenge's difficulty is whether a zero value represents a value below
365 the detection limit (rounded zero) or a truly absent taxon (essential zero). Here, we

366 employ zero replacement, which implies an assumption that all zero values represent
367 rounded zeros. New mixture models that explicitly allow for both essential and rounded
368 zeros (59), as well as more advanced methods of zero replacement (58, 60), may
369 enable us to handle zeros in a more sophisticated manner. Lastly, in regards to
370 informational loss caused by normalization, it is known that the number of counts
371 measured for a given taxon influences the precision with which we may estimate its
372 relative abundance in a sample (7, 32, 33). While our taxa weights are intended to
373 address this idea, a fully probabilistic model of counts would likely provide more
374 accurate error bounds for inference. We believe it would be possible to build such a
375 model in a Bayesian framework by viewing the observed counts as multinomial draws
376 from a point in the PhILR transformed space, as has been done for other log-ratio
377 based spaces (61).

378 Beyond refining the PhILR transform itself, future effort may also be directed
379 towards interpreting the transform's results at the single taxon level. Microbiota studies
380 frequently focus on individual taxa for tasks such as identifying specific bacteria that are
381 causal or biomarkers of disease. Log-ratio approaches can provide a compositionally
382 robust approach to identifying biomarkers based on changes in the relative abundance
383 of individual taxa. Due to the one-to-one correspondence between CLR coordinates and
384 individual taxa, the CLR transform has been used previously to build compositionally
385 robust models in terms of individual taxa (11, 19, 21). However, CLR transformed data
386 suffers from the drawback of a singular covariance matrix, which can make the
387 development of new models based on the CLR transform difficult (10). ILR transformed
388 data do not suffer this drawback (10) and moreover, can be analyzed at the single taxon

389 level. To do so, the inverse ILR transform can be applied to model results generated in
390 an ILR coordinate system, yielding analyses in terms of changes in the relative
391 abundance of individual taxa (10). The use of the inverse ILR transform in this manner
392 is well established (10, 62, 63) and the inverse transform is provided in the *Methods*
393 (34).

394 Despite these avenues for improvement, modification, or extension we believe
395 the PhILR transform already enables existing statistical methods to be applied to
396 metagenomic datasets, free from compositional artifacts and framed according to an
397 evolutionary perspective. We foresee the PhILR transform being used as a default
398 transformation prior to many microbiota analyses, particularly if a phylogenetic
399 perspective is desired. For example, the PhILR transform could be used in lieu of the
400 conventional log transform, which is often the default choice in microbiota analysis but
401 not robust to compositional effects. Substituting PhILR into existing bioinformatics
402 pipelines should often be seamless and we emphasize that all statistical tools applied to
403 PhILR transformed data in this study were used 'off-the-shelf' and without modification.
404 Importantly, such a substitution contrasts with the alternative approach for accounting
405 for compositional microbiota data, which is to modify existing statistical techniques.
406 Such modification is often challenging because many statistics were derived assuming
407 an unconstrained space with an orthonormal basis, not a constrained and over-
408 determined compositional space. Therefore, while select techniques have already been
409 adapted (e.g. distance measures that incorporate phylogenetic information (35) and
410 feature selection methods that handle compositional input (14, 15)), it is likely that
411 certain statistical goals, such as non-linear community forecasting or control system

412 modeling, may prove too complex for adapting to the compositional nature of microbiota
413 datasets. Finally, beyond microbiota surveys, we also recognize that compositional
414 metagenomics datasets are generated when studying the ecology of viral communities
415 (64) or clonal population structure in cancer (65-67). We expect the PhILR transform to
416 aid other arenas of biological research where variables are measured by relative
417 abundance and related by an evolutionary tree.

418

419 METHODS

420 *The ILR Transform*

421 A typical microbiome sample consists of measured counts c_j for taxa $j \in$
422 $\{1, \dots, D\}$. A standard operation is to take count data and transform it to relative
423 abundances. This operation is referred to as closure in compositional data analysis (5)
424 and is given by

$$x = \mathcal{C}[(c_1, \dots, c_D)] = \left(\frac{c_1}{\sum_j c_j}, \dots, \frac{c_D}{\sum_j c_j} \right)$$

425 where x represents a vector of relative abundances for the D taxa in the sample. We
426 can represent a binary phylogenetic tree of the D taxa using a sign matrix Θ as
427 introduced by Pawlowsky-Glahn and Egozcue (68) and shown in Figure 5. Each row of
428 the sign matrix indexes an internal node i of the tree and each column indexes a tip of
429 the tree. A given element in the sign matrix is ± 1 depending on which of the two clades
430 descending from i that tip is a part of and 0 if that tip is not a descendent of i . The
431 assignment of +1 versus -1 determines which clade is represented in the numerator
432 versus the denominator of the corresponding log-ratio (as described below).
433 Exchanging this assignment for a given balance switches which clade is represented in

434 the numerator versus the denominator of the log-ratio. Following Egozcue and
 435 Pawlowsky-Glahn (34), we represent the ILR coordinate (balance) associated with node
 436 i in terms of the shifted composition $\mathbf{y} = \mathbf{x}/\mathbf{p} = (x_1/p_1, \dots, x_D/p_D)$ as

$$y_i^* = \sqrt{\frac{n_i^+ n_i^-}{n_i^+ + n_i^-}} \log \frac{g_p(\mathbf{y}_i^+)}{g_p(\mathbf{y}_i^-)}. \quad (1)$$

437 Here, $g_p(\mathbf{y}_i^+)$ and $g_p(\mathbf{y}_i^-)$ represents the weighted geometric mean of the components
 438 of \mathbf{y} that represent tips in the $+1$ or -1 clade descendant from node i respectively. This
 439 weighted geometric mean is given by

$$g_p(\mathbf{y}_i^\pm) = \exp\left(\frac{\sum_{(\theta_{ij}=\pm 1)} p_j \log y_j}{\sum_{(\theta_{ij}=\pm 1)} p_j}\right) \quad (2)$$

440 where p_j is the weight assigned to taxa j . The term $\sqrt{n_i^+ n_i^- / n_i^+ + n_i^-}$ in equation 1 is the
 441 scaling term that ensures that the ILR basis element has unit length and the terms n_i^\pm
 442 are given by

$$n_i^\pm = \sum_{\theta_{ij}=\pm 1} p_j. \quad (3)$$

443 Note that when $\mathbf{p} = (1, \dots, 1)$, $\mathbf{y} = \mathbf{x}$, equation 1 represent the ILR transform as originally
 444 published (23), equation 2 represents the standard formula for geometric mean of a
 445 vector \mathbf{y} , and equation 3 represents the number of tips that descend from the $+1$ or -1
 446 clade descendant from node i . However, when $\mathbf{p} \neq (1, \dots, 1)$, these three equations
 447 represent a more generalized form of the ILR transform that allows weights to be
 448 assigned to taxa in the transformed space (34).

449 Following Egozcue and Pawlowsky-Glahn (34), we also note that the form of the
 450 generalized ILR (which we will denote ilr_p) transform can be rewritten in terms of a

451 generalized CLR transform (which we will denote clr_p). This formulation in terms of the
 452 generalized CLR transform can be more efficient to compute and allows the inverse of
 453 the transform to be easily described. We can define the generalized CLR transform as

$$454 \quad \text{clr}_p(\mathbf{y}) = \left(\log \frac{y_1}{g_p(y)}, \dots, \log \frac{y_D}{g_p(y)} \right).$$

455 The generalized ILR transform can then be written as

$$\mathbf{y}^* = \text{ilr}_p(\mathbf{y}) = \text{clr}_p(\mathbf{y}) \text{ diag}(\mathbf{p}) \Psi^T$$

456 with the ij th element of the matrix Ψ given by

$$\psi_{ij} = \begin{cases} +\frac{1}{n_i^+} \sqrt{\frac{n_i^+ n_i^-}{n_i^+ + n_i^-}} & \text{if } \theta_{ij} = +1 \\ -\frac{1}{n_i^-} \sqrt{\frac{n_i^+ n_i^-}{n_i^+ + n_i^-}} & \text{if } \theta_{ij} = -1 \\ 0 & \text{if } \theta_{ij} = 0. \end{cases}$$

457 With these components defined the inverse of generalized ILR transform can be written
 458 as $\mathcal{C}[\mathbf{y}] = \text{ilr}_p^{-1}(\mathbf{y}^*) = \mathcal{C}[\exp(\mathbf{y}^* \Psi)]$ and $\mathbf{x} = \mathcal{C}[\text{ilr}_p^{-1}(\mathbf{y}^*) \mathbf{p}]$.

459

460 **Soft thresholding through weighting taxa**

461 We make use of this generalized ILR transform to down weight the influence of
 462 taxa with many zero and near-zero counts since these are less reliable and therefore
 463 more variable (32). Our choice of taxa weights is a heuristic that combines two terms
 464 multiplicatively: a measure of the central tendency of counts, such as the mean or
 465 median of the raw counts for a taxon across the N samples in a dataset; and, the norm
 466 of the vector of relative abundances of a taxon across the N samples in a dataset. We
 467 add this vector norm term to weight taxa by their site-specificity. Preliminary studies

468 showed that the geometric mean of the counts (with a pseudocount added to avoid
469 skew from zero values) outperformed both the arithmetic mean and median as a
470 measure of central tendency for the counts (data not shown). Additionally, while both
471 the Euclidean norm and the Aitchison norm improved preliminary benchmark
472 performance compared to using the geometric mean alone, in one case (classification
473 using support vector machine on the global patterns dataset), the Euclidean norm
474 greatly outperformed the Aitchison norm (Supplementary file 1). Therefore, our chosen
475 taxa weighting scheme uses the geometric mean times the Euclidean norm:

476
$$p_j = \sqrt[N]{(c_{j1} + 1) \cdot \dots \cdot (c_{jN} + 1)} \cdot \|x_j\|.$$

477 Note that we add the subscript j to the right-hand side of the above equation to
478 emphasize that this is calculated with respect to a single taxon across the N samples in
479 a dataset. As intended, this scheme tended to assign smaller weights to taxa in our
480 benchmarks with more zero and near-zero counts (Figure 2-figure supplement 1).
481 Despite their heuristic nature, we found that our chosen weights provide performance
482 improvements over alternative weights (or the lack thereof) as measured by our
483 benchmark tasks (Supplementary File 1).

484 Our taxa weighting scheme supplements the use of pseudo-counts and
485 represents a soft-threshold on low abundance taxa. More generally, these taxa weights
486 represent a form of prior information regarding the importance of each taxon. We note
487 that if prior biological information suggests allowing specific taxa to influence the PhILR
488 transform more (or less) strongly, such a weighting could be achieved for taxon j by
489 increasing (or decreasing) p_j .

490

491 *Incorporating branch lengths*

492 Beyond utilizing the connectivity of the phylogenetic tree to dictate the
493 partitioning scheme for ILR balances, branch length information can be embedded into
494 the transformed space by linearly scaling ILR balances (y_i^*) by the distance between
495 neighboring clades. We call this scaling by phylogenetic distance ‘branch length
496 weighting’. Specifically, for each coordinate y_i^* , corresponding to node i we use the
497 transform

$$y_i^{*,blw} = y_i^* \cdot f(d_i^+, d_i^-)$$

498 where d_i^\pm represent the branch lengths of the two direct children of node i . When
499 $f(d_i^+, d_i^-) = 1$, the coordinates are not weighted by branch lengths. The form of this
500 transform was chosen so that the weights d_i^\pm , only influence the corresponding
501 coordinate ($y_i^{*,blw}$).

502 We also investigated the effect of using $f(d_i^+, d_i^-) = 1$, $f(d_i^+, d_i^-) = d_i^+ + d_i^-$, and
503 based on the results of Chen et al. (42), $f(d_i^+, d_i^-) = \sqrt{d_i^+ + d_i^-}$ on benchmark
504 performance. When coupled with the taxa weights specified above, the square root of
505 the summed distances had the highest rank in 9 of the 12 supervised classification
506 tasks and 2 of the 3 distance based tasks (Supplementary file 1). Based on these
507 results, except for our analysis of balance variance versus phylogenetic depth (see
508 below), the square root of the summed distances was used throughout our analyses.

509

510 *Implementation*

511 The PhILR transform, as well as the incorporation of branch length and taxa
512 weightings has been implemented in the R programming language as the package *philr*
513 available at <https://bioconductor.org/packages/philtr/>.

514

515 ***Datasets and Preprocessing***

516 All data preprocessing was done in the R programming language using the
517 *phyloseq* package for analysis of microbiome census data (69) as well as the *ape* (70)
518 and *phangorn* (71) packages for analysis of phylogenetic trees.

519 ***Data Acquisition***

520 We chose to use previously published OTU tables, taxonomic classifications, and
521 phylogenies as the starting point for our analyses. The Human Microbiome Project
522 (HMP) dataset was obtained from the QIIME Community Profiling Pipeline applied to
523 high-quality reads from the v3-5 region, available at <http://hmpdacc.org/HMQCP/>. The
524 Global Patterns dataset was originally published in Caporaso, *et al.* (1) and is provided
525 with the *phyloseq* R package (69). The Costello Skin Sites dataset (CSS) is a subset of
526 the dataset collected by Costello *et al.* (39) featuring only the samples from skin sites.
527 This skin subset was introduced as a benchmark for supervised machine learning by
528 Knights *et al.* (40) and can be obtained from <http://knightslab.org/data>.

529 ***OTU Table Preprocessing***

530 To accord with general practice, we performed a minimal level of OTU table
531 filtering for all datasets used in benchmarks and analyses. Due to differences in
532 sequencing depth, sequencing methodology, and the number and diversity of samples
533 between datasets, filtering thresholds were set independently for each dataset. For the

534 HMP dataset, we initially removed samples with fewer than 1000 counts to mimic prior
535 analyses (41). We additionally removed OTUs that were not seen with more than 3
536 counts in at least 1% of samples. Preprocessing of the Global Patterns OTU table
537 followed the methods outlined in McMurdie and Holmes (69). Specifically, OTUs that
538 were not seen with more than 3 counts in at least 20% of samples were removed, the
539 sequencing depth of each sample was standardized to the abundance of the median
540 sampling depth, and finally OTUs with a coefficient of variation ≤ 3.0 were removed. The
541 CSS dataset had lower sequencing depth than the other two datasets; we chose to filter
542 OTUs that were not seen with greater than 10 counts across the skin samples. The
543 PhILR transform, and more generally our benchmarking results in Figure 2b and 2C,
544 were robust to varying our filtering strategies (Supplementary File 2).

545 *Preprocessing Phylogenies*

546 For each dataset, the phylogeny was pruned to include only those taxa remaining
547 after OTU table preprocessing. Except for the Global Patterns dataset, which was
548 already rooted, we chose to root phylogenies by manually specifying an outgroup. For
549 the HMP dataset the phylum Euryarchaeota was chosen as an outgroup. For the CSS
550 dataset, the tree was rooted with OTU 12871 (from phylum Plantomycetes) as the
551 outgroup. For all three phylogenies, any multichotomies were resolved using with the
552 function multi2di from the *ape* package which replaces multichotomies with a series of
553 dichotomies with one (or several) branch(es) of length zero.

554 *Zero replacement and Normalization*

555 A pseudocount of 1 was added prior to PhILR transformation to avoid taking log-
556 ratios with zero counts. We found that our benchmarking results were robust to
557 changing the value of this pseudocount from 1 to 2, 3, or 10 (Supplementary file 1).

558 *Grouping Sampling Sites*

559 To simplify subsequent analyses, HMP samples from the left and right
560 retroauricular crease and samples from the left and right antecubital fossa were
561 grouped together, respectively, as preliminary PERMANOVA analysis suggested that
562 these sites were indistinguishable (data not shown).

563

564 **Benchmarking**

565 *Distance/Dissimilarity Based Analysis*

566 Distance between samples in PhILR transformed space was calculated using
567 Euclidean distance. All other distance measures were calculated using *phyloseq* on the
568 preprocessed data without adding a pseudocount. Principle coordinate analysis was
569 performed for visualization using *phyloseq*. PERMANOVA was performed using the
570 function *adonis* from the R package *vegan* (v2.3.4). The R^2 value from the fitted model
571 was taken as a performance metric. Standard errors were calculated using bootstrap
572 resampling with 100 samples each. Differences between the performance of Euclidean
573 distance in PhILR transformed space and that of each other distance or dissimilarity
574 measure on a given task was tested using two-sided t-tests and multiple hypothesis
575 testing was accounted for using FDR correction.

576 *Supervised Classification*

577 The performance, as measured by classification accuracy, of PhILR transformed
578 data was compared against data preprocessed using one of two standard strategies for
579 normalizing sequencing depth: the preprocessed data was transformed to relative
580 abundances (e.g., each sample was normalized to a constant sum of 1; *raw*); or, a
581 pseudocount of 1 was added, the data was transformed to relative abundances, and
582 finally the relative abundances were log-transformed (*log*).

583 All supervised learning was implemented in Python using the following libraries:
584 *Scikit-learn* (v0.17.1), *numpy* (v1.11.0) and *pandas* (v0.17.1). Four classifiers were
585 used: penalized logistic regression, support vector classification with RBF kernel,
586 random forest classification, and k-nearest-neighbors classification. Each classification
587 task was evaluated using the mean and variance of the test accuracy over 10
588 randomized test/train (30/70) splits which preserved the percentage of samples from
589 each class at each split. For each classifier, for each split, the following parameters
590 were set using cross-validation on the training set. Logistic regression and Support
591 Vector Classification: the ‘C’ parameter was allowed to vary between 10^{-3} to 10^3 and
592 multi-class classification was handled with a one-vs-all loss. In addition, for logistic
593 regression the penalty was allowed to be either l_1 or l_2 . K-nearest-neighbors
594 classification: the ‘weights’ argument was set to ‘distance’. Random forest classification:
595 each forest contained 30 trees and the ‘max_features’ argument was allowed to vary
596 between 0.1 and 1. All other parameters were set to default values. Due to the small
597 size of the Global Patterns dataset, the supervised classification task was simplified to
598 distinguishing human vs. non-human samples. Differences between each methods’

599 accuracy in each task was tested using two-sided t-tests and multiple hypothesis testing
600 was accounted for using FDR correction.

601

602 ***Identifying balances that distinguish sites***

603 To identify a sparse set of balances that distinguish sampling sites while
604 accounting for the dependencies between nested balances, we fit a multinomial
605 regression model with a grouped l_1 penalty using the R package *glmnet* (*v2.0.5*). The
606 penalization term lambda was set by visually inspecting model outputs for clear body
607 site separation (lambda=0.1198). This resulted in 35 balances with non-zero regression
608 coefficients. Phylogenetic tree visualization was done using the R package *ggtree* (72).

609

610 ***Variance and Depth*** To reduce the likelihood that our analysis of balance
611 variance and phylogenetic depth was affected by statistical artifact, we modified our
612 PhILR transform in several ways. First, we omitted branch length weights (*i.e.*, we set
613 $f(d_i^+, d_i^-) = 1$) as these may vary non-randomly as a function of phylogenetic depth.
614 Second, we also anticipated that any zero replacement method would likely lead to
615 lower variance measurements, which could have greater effects on balances closer to
616 the tips of the tree. We therefore omitted taxa weights and zero replacement; we
617 instead used stricter hard filtering thresholds and calculated balance values based on
618 non-zero counts. In practice, we used the following filtering thresholds for each body
619 site, taxa present in less than 20% of samples from that site were excluded and
620 subsequently samples that had less than 50 total counts were excluded. To calculate
621 balance values based on non-zero counts we retained balances that met the following

622 criteria: the term $g_p(\mathbf{y}_i^+)/g_p(\mathbf{y}_i^-)$ had non-zero counts from some taxa within the
623 subcomposition \mathbf{y}_i^+ (formed by the taxa that descend from the +1 clade of node i) and
624 some other taxa within the subcomposition \mathbf{y}_i^- (formed by the taxa that descend from
625 the -1 clade of node i) in at least 40 samples from that body site. We believe these two
626 modifications to PhILR resulted in a more conservative analysis for of balance variance
627 versus phylogenetic depth but are likely not optimal in other situations.

628 To investigate the overall relationship between balance variance and
629 phylogenetic depth we used linear regression. A balance's depth in the tree was
630 calculated as its mean phylogenetic distance to its descendant tips (d). For a given
631 body site the following model was fit:

$$\log \text{var}(y^*) = \beta \log d + \alpha$$

632 where d represents mean distance from a balance to its descendant tips. We then set
633 out to test the null hypothesis that $\beta = 0$, or that the variance of the log-ratio between
634 two clades was invariant to the distance of the two clades from their most recent
635 common ancestor. For each site, a null distribution for β was constructed by
636 permutations of the tip labels of the phylogenetic tree. For each permutation of the
637 labels, the resultant tree was used to transform the data and β was estimated. We
638 chose this permutation scheme to ensure that the increasing variance we saw with
639 increasing proximity of a balance to the root was not because deeper balances had
640 more descendant tips, an artifact of variance scaling with mean abundance, or due to
641 bias introduced due to our handling of zeros. Furthermore, for each body site, we found
642 the null distribution for β was symmetric about $\beta = 0$ which further supports that balance
643 variance depends on phylogenetic depth through a biological mechanism and not

644 through a statistical artifact (Figure 4-figure supplement 3). Two tailed p-values were
645 calculated for β based on 20000 samples from each site's respective null distribution.
646 FDR correction was applied to account for multiple hypothesis testing between body
647 sites.

648 To visualize local trends in the relationship between balance variance and
649 phylogenetic depth, a LOESS regression was fit independently for each body site. This
650 was done using the function *geom_smooth* from the R package *ggplot2* (v2.1.0) with
651 default parameters.

652 The data and code needed to reproduce our analysis of balance variance versus
653 phylogenetic depth is provided in Figure 4 – source data and Figure 4 – source code
654 respectively.

655

656 ***Integrating Taxonomic Information***

657 Taxonomy was assigned to OTUs in the HMP dataset using the
658 *assign_taxonomy.py* script from *Qiime* (v1.9.1) to call *uclust* (v1.2.22) with default
659 parameters using the representative OTU sequences obtained as described above.
660 Taxonomic identifiers were assigned to the two descendant clades of a given balance
661 separately using a simple voting scheme and combined into a single name for that
662 balance. The voting scheme occurs as follows: (1) for a given clade, the entire
663 taxonomy table was subset to only contain the OTUs that were present in that clade (2)
664 starting at the finest taxonomic rank the subset taxonomy table was checked to see if
665 any species identifier represented $\geq 95\%$ of the table entries at that taxonomic rank, if so

666 that identifier was taken as the taxonomic label for the clade (3) if no consensus
667 identifier was found, the table was checked at the next most-specific taxonomic rank.

668 Median phylogenetic depths for each taxonomic rank were estimated by first
669 decorating a phylogenetic tree with taxonomy information using *tax2tree* (v1.0) (73). For
670 a given taxonomic rank the mean distance to tips was calculated for each internal node
671 possessing a label that ended in that rank. The median of these distances was used to
672 display an estimate of the phylogenetic depth of that given rank. This calculation of
673 median phylogenetic depth of different taxonomic ranks was done separately for each
674 body site.

675 The data and code needed to reproduce the taxonomic assignment and
676 estimation of median phylogenetic depths for each taxonomic rank is included in Figure
677 4 – source data and Figure 4 – source code respectively.

678

679
680

681 **ACKNOWLEDGEMENTS:**

682 We thank Rachel Silverman, Aspen Reese, Firas Midani, Heather Durand, Jesse
683 Shapiro, Jonathan Friedman, Simon Levin, and Susan Holmes for their helpful
684 comments, Dan Knights for providing us with the CSS dataset, as well as Klaus Schliep
685 and Liam Revell for their insight into manipulation of phylogenetic trees in the R
686 programming language.

687

688 **COMPETING INTERESTS:**

689 The authors declare that they have no competing financial, professional, or personal
690 interests that might have influenced the performance or presentation of the work
691 described in this manuscript.

692

693 **REFERENCES**

- 694 1. Caporaso JG, et al. (2011) Global patterns of 16S rRNA diversity at a depth of
695 millions of sequences per sample. *P Natl Acad Sci USA* 108 Suppl 1:4516-4522.
- 696 2. Waldor MK, et al. (2015) Where next for microbiome research? *PLoS Biol*
697 13(1):e1002050.
- 698 3. Jackson DA (1997) Compositional data in community ecology: The paradigm or
699 peril of proportions? *Ecology* 78(3):929-940.
- 700 4. Friedman J & Alm EJ (2012) Inferring correlation networks from genomic survey
701 data. *PLoS Comput Biol* 8(9):e1002687.
- 702 5. Aitchison J (1986) *The statistical analysis of compositional data* (Chapman and
703 Hall, London ; New York).
- 704 6. Lovell D, Müller W, Taylor J, Zwart A, & Helliwell C (2011) Proportions,
705 percentages, ppm: do the molecular biosciences treat compositional data right.
706 *Compositional Data Analysis: Theory and Applications*, eds Pawlowsky-Glahn V
707 & Buccianti A (John Wiley & Sons, Ltd.), pp 193-207.
- 708 7. Gloor GB, Macklaim JM, Vu M, & Fernandes AD (2016) Compositional
709 uncertainty should not be ignored in high-throughput sequencing data analysis.
710 *Austrian Journal of Statistics* 45(4):73.
- 711 8. Li HZ (2015) Microbiome, Metagenomics, and High-Dimensional Compositional
712 Data Analysis. *Annual Review of Statistics and Its Application*, Vol 2 2(1):73-94.
- 713 9. Tsilimigras MC & Fodor AA (2016) Compositional data analysis of the
714 microbiome: fundamentals, tools, and challenges. *Ann Epidemiol* 26(5):330-335.
- 715 10. Pawlowsky-Glahn V, Egozcue JJ, & Tolosana-Delgado R (2015) *Modeling and
716 analysis of compositional data* (John Wiley & Sons, Ltd).
- 717 11. Mandal S, et al. (2015) Analysis of composition of microbiomes: a novel method
718 for studying microbial composition. *Microb Ecol Health Dis* 26:27663.
- 719 12. Fang H, Huang C, Zhao H, & Deng M (2015) CCLasso: correlation inference for
720 compositional data through Lasso. *Bioinformatics* 31(19):3172-3180.
- 721 13. La Rosa PS, et al. (2012) Hypothesis testing and power calculations for
722 taxonomic-based human microbiome data. *PLoS one* 7(12):e52078-e52078.
- 723 14. Chen J & Li H (2013) Variable Selection for Sparse Dirichlet-Multinomial
724 Regression with an Application to Microbiome Data Analysis. *Ann Appl Stat*
725 7(1):418-442.
- 726 15. Lin W, Shi PX, Feng R, & Li H (2014) Variable selection in regression with
727 compositional covariates. *Biometrika* 101(4):785-797.
- 728 16. Paulson JN, Stine OC, Bravo HC, & Pop M (2013) Differential abundance
729 analysis for microbial marker-gene surveys. *Nat Methods* 10(12):1200-1202.
- 730 17. Anders S & Huber W (2010) Differential expression analysis for sequence count
731 data. *Genome Biol* 11(10):R106.
- 732 18. Bacon-Shone J (2011) A Short History of Compositional Data Analysis.
733 *Compositional Data Analysis*, eds Pawlowsky-Glahn V & Buccianti A (John Wiley
734 & Sons, Ltd), pp 1-11.
- 735 19. Kurtz ZD, et al. (2015) Sparse and compositionally robust inference of microbial
736 ecological networks. *PLoS Comput Biol* 11(5):e1004226.
- 737 20. Lee SC, et al. (2014) Helminth colonization is associated with increased diversity
738 of the gut microbiota. *PLoS Negl Trop Dis* 8(5):e2880.

- 739 21. Fernandes AD, et al. (2014) Unifying the analysis of high-throughput sequencing
740 datasets: characterizing RNA-seq, 16S rRNA gene sequencing and selective
741 growth experiments by compositional data analysis. *Microbiome* 2:15.
- 742 22. Gloor GB, Wu JR, Pawlowsky-Glahn V, & Egozcue JJ (2016) It's all relative:
743 analyzing microbiome data as compositions. *Ann Epidemiol* 26(5):322-329.
- 744 23. Egozcue JJ, Pawlowsky-Glahn V, Mateu-Figueras G, & Barcelo-Vidal C (2003)
745 Isometric logratio transformations for compositional data analysis. *Mathematical
746 Geology* 35(3):279-300.
- 747 24. Egozcue JJ & Pawlowsky-Glahn V (2005) Groups of parts and their balances in
748 compositional data analysis. *Mathematical Geology* 37(7):795-828.
- 749 25. Finucane MM, Sharpton TJ, Laurent TJ, & Pollard KS (2014) A taxonomic
750 signature of obesity in the microbiome? Getting to the guts of the matter. *PLoS
751 One* 9(1):e84689.
- 752 26. Le Cao KA, et al. (2016) MixMC: A Multivariate Statistical Framework to Gain
753 Insight into Microbial Communities. *PLoS One* 11(8):e0160169.
- 754 27. Morton JT, et al. (2017) Balance Trees Reveal Microbial Niche Differentiation.
755 *mSystems* 2(1).
- 756 28. Hunt DE, et al. (2008) Resource partitioning and sympatric differentiation among
757 closely related bacterioplankton. *Science* 320(5879):1081-1085.
- 758 29. De Filippo C, et al. (2010) Impact of diet in shaping gut microbiota revealed by a
759 comparative study in children from Europe and rural Africa. *P Natl Acad Sci USA*
760 107(33):14691-14696.
- 761 30. Wu GD, et al. (2011) Linking long-term dietary patterns with gut microbial
762 enterotypes. *Science* 334(6052):105-108.
- 763 31. Yatsunenko T, et al. (2012) Human gut microbiome viewed across age and
764 geography. *Nature* 486(7402):222-227.
- 765 32. Good IJ (1956) On the Estimation of Small Frequencies in Contingency-Tables.
766 *Journal of the Royal Statistical Society Series B-Statistical Methodology*
767 18(1):113-124.
- 768 33. McMurdie PJ & Holmes S (2014) Waste not, want not: why rarefying microbiome
769 data is inadmissible. *PLoS Comput Biol* 10(4):e1003531.
- 770 34. Egozcue JJ & Pawlowsky-Glahn V (2016) Changing the Reference Measure in
771 the Simplex and its Weightings Effects. *Austrian Journal of Statistics* 45(4):25-44.
- 772 35. Lozupone C & Knight R (2005) UniFrac: a new phylogenetic method for
773 comparing microbial communities. *Appl Environ Microbiol* 71(12):8228-8235.
- 774 36. Fukuyama J, McMurdie PJ, Dethlefsen L, Relman DA, & Holmes S (2012)
775 Comparisons of distance methods for combining covariates and abundances in
776 microbiome studies. *Pacific Symposium on Biocomputing.*, (NIH Public Access),
777 p 213.
- 778 37. Purdom E (2011) Analysis of a Data Matrix and a Graph: Metagenomic Data and
779 the Phylogenetic Tree. *Annals of Applied Statistics* 5(4):2326-2358.
- 780 38. Martiny JB, Jones SE, Lennon JT, & Martiny AC (2015) Microbiomes in light of
781 traits: A phylogenetic perspective. *Science* 350(6261):aac9323.
- 782 39. Costello EK, et al. (2009) Bacterial community variation in human body habitats
783 across space and time. *Science* 326(5960):1694-1697.

- 784 40. Knights D, Costello EK, & Knight R (2011) Supervised classification of human
785 microbiota. *FEMS Microbiol Rev* 35(2):343-359.
- 786 41. Consortium HMP (2012) Structure, function and diversity of the healthy human
787 microbiome. *Nature* 486(7402):207-214.
- 788 42. Chen J, et al. (2012) Associating microbiome composition with environmental
789 covariates using generalized UniFrac distances. *Bioinformatics* 28(16):2106-
790 2113.
- 791 43. Lovell D, Pawlowsky-Glahn V, Egoscue JJ, Marguerat S, & Bahler J (2015)
792 Proportionality: a valid alternative to correlation for relative data. *PLoS Comput
793 Biol* 11(3):e1004075.
- 794 44. Matsen FA (2015) Phylogenetics and the human microbiome. *Syst Biol*
795 64(1):e26-41.
- 796 45. Moeller AH, et al. (2016) Cospeciation of gut microbiota with hominids. *Science*
797 353(6297):380-382.
- 798 46. Ley RE, et al. (2008) Evolution of mammals and their gut microbes. *Science*
799 320(5883):1647-1651.
- 800 47. Smillie CS, et al. (2011) Ecology drives a global network of gene exchange
801 connecting the human microbiome. *Nature* 480(7376):241-244.
- 802 48. Rakoff-Nahoum S, Foster KR, & Comstock LE (2016) The evolution of
803 cooperation within the gut microbiota. *Nature* 533(7602):255-259.
- 804 49. Blaser MJ & Falkow S (2009) What are the consequences of the disappearing
805 human microbiota? *Nat Rev Microbiol* 7(12):887-894.
- 806 50. Aas JA, Paster BJ, Stokes LN, Olsen I, & Dewhirst FE (2005) Defining the
807 normal bacterial flora of the oral cavity. *J Clin Microbiol* 43(11):5721-5732.
- 808 51. Mager DL, Ximenez-Fyvie LA, Haffajee AD, & Socransky SS (2003) Distribution
809 of selected bacterial species on intraoral surfaces. *J Clin Periodontol* 30(7):644-
810 654.
- 811 52. Grice EA & Segre JA (2011) The skin microbiome. *Nat Rev Microbiol* 9(4):244-
812 253.
- 813 53. Janda JM & Abbott SL (2007) 16S rRNA gene sequencing for bacterial
814 identification in the diagnostic laboratory: pluses, perils, and pitfalls. *J Clin
815 Microbiol* 45(9):2761-2764.
- 816 54. Vetrovsky T & Baldrian P (2013) The variability of the 16S rRNA gene in bacterial
817 genomes and its consequences for bacterial community analyses. *PLoS One*
818 8(2):e57923.
- 819 55. Zaneveld JR, Lozupone C, Gordon JI, & Knight R (2010) Ribosomal RNA
820 diversity predicts genome diversity in gut bacteria and their relatives. *Nucleic
821 Acids Res* 38(12):3869-3879.
- 822 56. Levy R & Borenstein E (2013) Metabolic modeling of species interaction in the
823 human microbiome elucidates community-level assembly rules. *P Natl Acad Sci
824 USA* 110(31):12804-12809.
- 825 57. Faust K, et al. (2012) Microbial co-occurrence relationships in the Human
826 Microbiome. *PLoS Computational Biology* 8(7).
- 827 58. Martín-Fernandez JA, Palarea-Albaladejo J, & Olea RA (2011) Dealing with
828 zeros. *Compositional data analysis: Theory and applications*, eds Pawlowsky-
829 Glahn V & Buccianti A), pp 43-58.

- 830 59. Bear J & Billheimer D (2016) A Logistic Normal Mixture Model for Compositional
831 Data Allowing Essential Zeros. *Austrian Journal of Statistics* 45(4):3-23.
- 832 60. Martín-Fernández J-A, Hron K, Templ M, Filzmoser P, & Palarea-Albaladejo J
833 (2015) Bayesian-multiplicative treatment of count zeros in compositional data
834 sets. *Statistical Modelling* 15(2):134-158.
- 835 61. Billheimer D, Guttorp P, & Fagan WF (2001) Statistical interpretation of species
836 composition. *Journal of the American Statistical Association* 96(456):1205-1214.
- 837 62. Boogaart Gvd & ebrary Inc. (2013) Analyzing compositional data with R. in
838 *UseR!* (Springer,, New York), pp 1 online resource (xv, 258 p.
- 839 63. Pawlowsky-Glahn V & Buccianti A (2011) Compositional data analysis : theory
840 and applications. (Wiley,, Chichester, West Sussex ; Hoboken, N.J.), pp xxi, 378
841 p.
- 842 64. Culley AI, Lang AS, & Suttle CA (2006) Metagenomic analysis of coastal RNA
843 virus communities. *Science* 312(5781):1795-1798.
- 844 65. Britanova OV, et al. (2014) Age-related decrease in TCR repertoire diversity
845 measured with deep and normalized sequence profiling. *J Immunol* 192(6):2689-
846 2698.
- 847 66. Yuan K, Sakoparnig T, Markowetz F, & Beerenwinkel N (2015) BitPhylogeny: a
848 probabilistic framework for reconstructing intra-tumor phylogenies. *Genome Biol*
849 16:36.
- 850 67. Roth A, et al. (2014) PyClone: statistical inference of clonal population structure
851 in cancer. *Nat Methods* 11(4):396-398.
- 852 68. Pawlowsky-Glahn V & Egozcue JJ (2011) Exploring Compositional Data with the
853 CoDa-Dendrogram. *Austrian Journal of Statistics* 40(1&2):103-113.
- 854 69. McMurdie PJ & Holmes S (2013) phyloseq: an R package for reproducible
855 interactive analysis and graphics of microbiome census data. *PLoS One*
856 8(4):e61217.
- 857 70. Paradis E, Claude J, & Strimmer K (2004) APE: Analyses of Phylogenetics and
858 Evolution in R language. *Bioinformatics* 20(2):289-290.
- 859 71. Schliep KP (2011) phangorn: phylogenetic analysis in R. *Bioinformatics*
860 27(4):592-593.
- 861 72. Yu G, Smith DK, Zhu H, Guan Y, & Lam TTY (2016) ggtree: an R package for
862 visualization and annotation of phylogenetic trees with their covariates and other
863 associated data. *Methods in Ecology and Evolution*.
- 864 73. McDonald D, et al. (2012) An improved Greengenes taxonomy with explicit ranks
865 for ecological and evolutionary analyses of bacteria and archaea. *ISME J*
866 6(3):610-618.

869 **Figure 1. PhILR uses an evolutionary tree to transform microbiota data into an unconstrained**
870 **coordinate system.** (A) Two hypothetical bacterial communities share identical absolute numbers of
871 *Lactobacillus*, and *Ruminococcus* bacteria; they differ only in the absolute abundance of *Bacteroides*
872 which is higher in community A (red circle) compared to community B (blue diamond). (B) A ternary plot
873 depicts proportional data typically analyzed in a sequencing-based microbiota survey. Note that viewed in
874 terms of proportions the space is constrained and the axes are not Cartesian. As a result, all three genera
875 have changed in relative abundance between the two communities. (C) Schematic of the PhILR transform
876 based on a phylogenetic sequential binary partition. The PhILR coordinates can be viewed as ‘balances’
877 between the weights (relative abundances) of the two subclades of a given internal node. In community
878 B, the greater abundance of *Bacteroides* tips the balance y_1^* to the right. (D) The PhILR transform can be
879 viewed as a new coordinate system (grey dashed lines) in the proportional data space. (E) The data
880 transformed to the PhILR space. Note that in contrast to the raw proportional data (B), the PhILR space
881 only shows a change in the variable associated with *Bacteroides*.

882

883 **Figure 2. Performance of standard statistical models on PhILR transformed microbiota data.**
884 Benchmarks were performed using three datasets: Costello Skin Sites (CSS), Global Patterns (GP),
885 Human Microbiome Project (HMP) (a summary of these datasets after preprocessing is shown in
886 Supplementary file 1 and Figure 2-figure supplement 1). (A) Sample distance visualized using principal
887 coordinate analysis (PCoA) of Euclidean distances computed in PhILR coordinate system. A comparison
888 to PCoAs calculated with other distance measures is shown in Figure 2-figure supplement 2. (B) Sample
889 distance (or dissimilarity) was computed by a range of statistics. PERMANOVA R^2 values, which
890 represent how well sample identity explained the variability in sample pairwise distances, were used as a
891 performance metric. Distances in the PhILR transformed space were calculated using Euclidean distance.
892 Distances between samples on raw relative abundance data were computed using Weighted and
893 Unweighted UniFrac (WUnifrac and Unifrac, respectively), Bray-Curtis, Binary Jaccard, and Euclidean
894 distance. Error bars represent standard error measurements from 100 bootstrap replicates and (*)
895 denotes a p-value of ≤ 0.01 after FDR correction of pairwise tests against PhILR. (C) Accuracy of
896 supervised classification methods tested on benchmark datasets. Error bars represent standard error
897 measurements from 10 test/train splits and (*) denotes a p-value of ≤ 0.01 after FDR correction of all
898 pairwise tests.

899

900

901 **Figure 2-figure supplement 1. Taxa weighting scheme tends to assign smaller weights to taxa with**
902 **more zero and near zero counts.** The weight of a given taxon is calculated as the geometric mean of its
903 counts across all samples times the Euclidean norm of its relative abundance across all samples in the
904 dataset (*Methods*). Data are plotted on a log scale with a pseudocount of 1 added to ease visualization.
905

906 **Figure 2-figure supplement 2. Principal coordinate analyses using different measures of**
907 **community distance or dissimilarity.**

908
909 **Figure 3. Balances distinguishing human microbiota by body site.** Sparse logistic regression was
910 used to identify balances that best separated the different sampling sites (full list of balances provided in
911 Figure 3-figure supplement 1). **(A)** Each balance is represented on the tree as a broken grey bar. The left
912 portion of the bar identifies the clade in the denominator of the log-ratio, and the right portion identifies the
913 clade in the numerator of the log-ratio. The branch leading from the Firmicutes to the Bacteroidetes has
914 been rescaled to facilitate visualization. **(B-F)** The distribution of balance values across body sites.
915 Vertical lines indicate median values, boxes represent interquartile ranges (IQR) and whiskers extend to
916 1.5 IQR on either side of the median. Balances between: **(B)** the phyla Actinobacteria and Fusobacteria
917 versus the phyla Bacteroidetes, Firmicutes, and Proteobacteria distinguish stool and oral sites from skin
918 sites; **(C)** *Prevotella spp.* and *Bacteroides spp.* distinguish stool from oral sites; **(D)** *Corynebacterium spp.*
919 distinguish skin and oral sites; **(E)** *Streptococcus spp.* distinguish oral sites; and, **(F)** *Actinomyces spp.*
920 distinguish oral plaques from other oral sites. (†) Includes Bacteroidetes, Firmicutes, Alpha-, Beta-, and
921 Gamma-proteobacteria. (‡) Includes Actinobacteria, Fusobacteria, Epsilon-proteobacteria, Spirochaetes,
922 and Verrucomicrobia.
923

924 **Figure 3-figure supplement 1. Balances found to distinguish human body sites by sparse logistic**
925 **regression.** Taxa are listed in balance *numerator/denominator* format.
926

927 **Figure 4. Neighboring clades covary less with increasing phylogenetic depth.** The variance of
928 balance values captures the degree to which neighboring clades covary, with smaller balance variances
929 representing sister clades that covary more strongly (Figure 4-figure supplement 1). **(A-C)** Balance
930 variances were computed among samples from stool **(A)**, buccal mucosa **(B)**, and the mid-vagina **(C)**.
931 Red branches indicate small balance variance and blue branches indicate high balance variance.
932 Balances 1-6 are individually tracked in panels **(D-L)**. **(D-F)** Balance variances within each body site
933 increased linearly with increasing phylogenetic depth on a log-scale (blue line; p<0.01, permutation test

934 with FDR correction; *Methods*). Significant trends are seen across all other body sites (Figure 4-figure
935 supplement 2 and 3). Non-parametric LOESS regression (green line and corresponding 95% confidence
936 interval) reveals an inflection point in the relation between phylogenetic depth and balance variance. This
937 inflection point appears below the estimated species level ('s' dotted line; the median depth beyond which
938 balances no longer involve leaves sharing the same species assignment; *Methods*). (**G-L**) Examples of
939 balances with high and low variance from panels (**A-F**). Low balance variances (**H, J, L**) reflect a linear
940 relationship between the geometric means of sister clades abundances. High balance variances (**G, I, K**)
941 reflect either unlinked or exclusionary dynamics between the geometric means of sister clades
942 abundances.

943

944 **Figure 4-figure supplement 1. Balances with high and low variance.** Shown are the 10 highest and
945 the 10 lowest variance balances for the Stool body site in the HMP dataset. Panels 1-10: Low balance
946 variances reflect a linear relationship (through the origin) between levels of the two clades that descend
947 from a given balance. Panels 11-20: High balances variances lack this linear relationship.

948

949 **Figure 4-figure supplement 2. Neighboring clades covary less with increasing phylogenetic depth.**
950 Shown here are all HMP body sites. Dotted lines indicate median species boundaries (see *Methods*).

951

952 **Figure 4-figure supplement 3. The null distribution for β .** β is the slope in the linear regression
953 between balance variance and phylogenetic depth in log space (regressions shown in Figure 4-figure
954 supplement 2). To form null distributions for β , tip labels on the phylogeny were shuffled ($n=20,000$) and
955 balance values re-calculated. Distributions for β are symmetric about 0 for each body site, suggesting
956

957 **Figure 5. Sign matrix representation of a phylogenetic tree.** A binary tree (**Left**) can be represented
958 by a sign matrix (**Right**) denoted Θ .

959

960

961

962

963 **SUPPLEMENTARY FILES**

964

965 **Supplementary file 1.** Extended benchmarking results including the effects of changing
966 pseudocount and different choices for taxa and branch length weights. This file also
967 contains summary information for the benchmark datasets after preprocessing.

968

969 **Supplementary file 2.** Extended benchmarking results regarding the sensitivity of
970 methods to different OTU filtering schemes.

971

972 **Figure 2 – source data.** Source data for Figure 2b and 2c as well as FDR corrected p-
973 values from tests.

974

975 **Figure 4 – source code.** Source code for Figure 4 and associated supplements.

976

977 **Figure 4 – source data.** FDR corrected p-values from permutation tests.

978

979

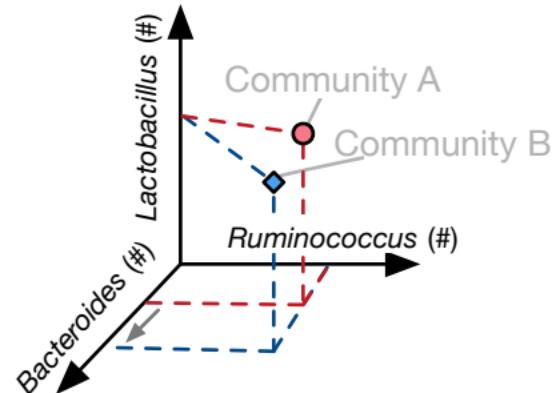
980

981

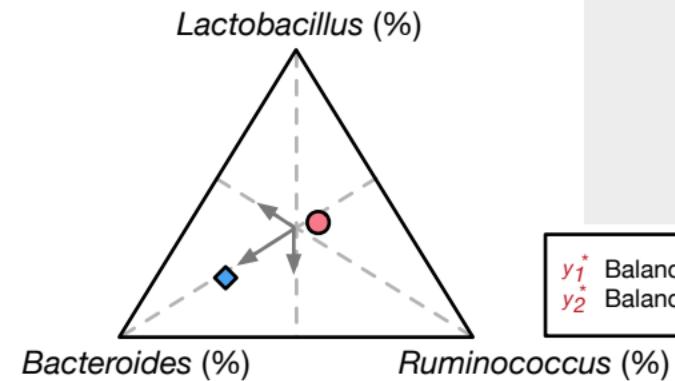
982

983

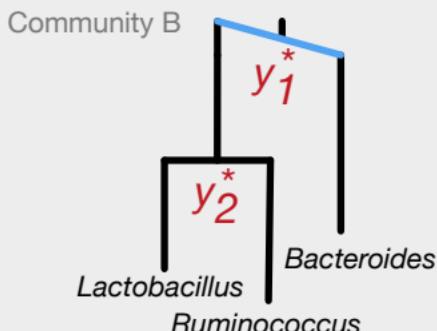
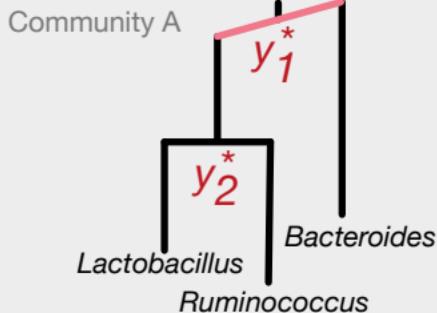
A Unobserved Absolute Abundances



B Observed Compositions

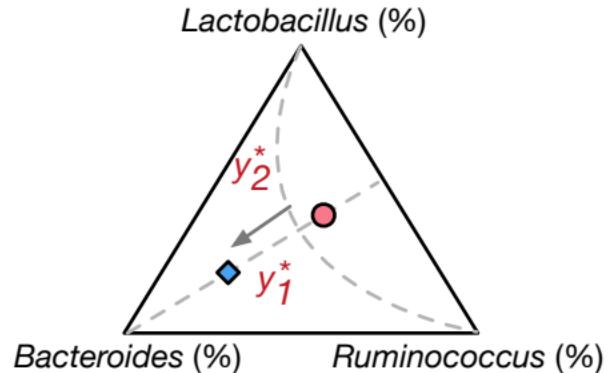


C Balances Depicted on Phylogenetic Tree

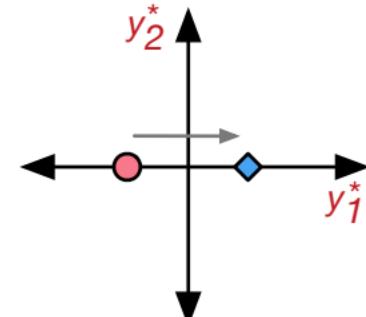


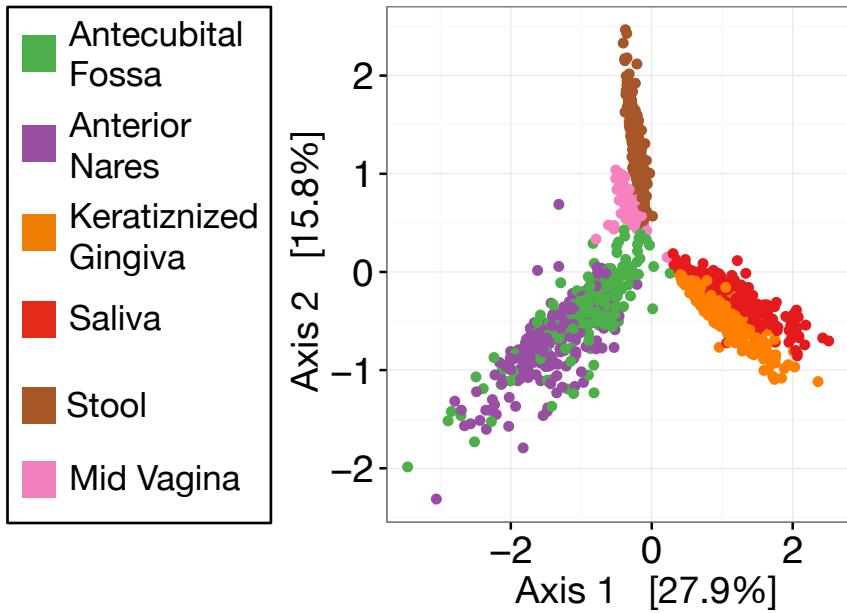
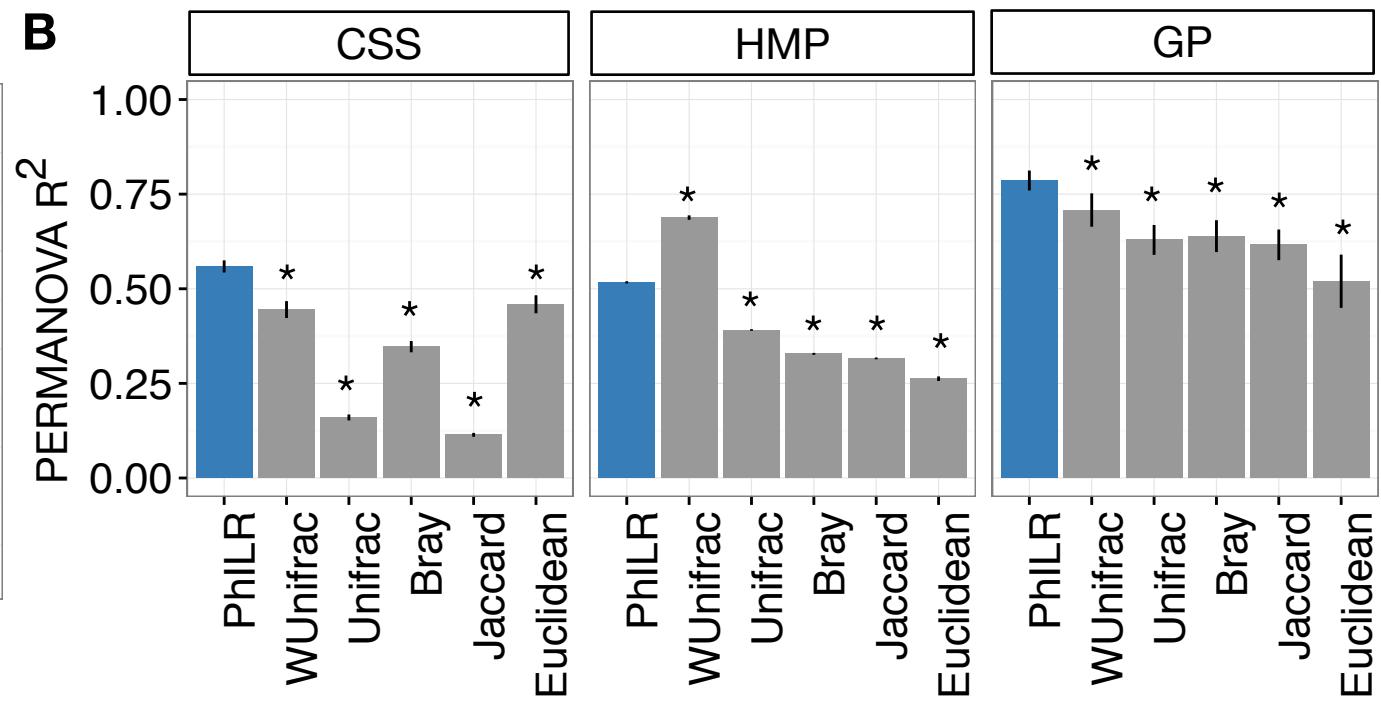
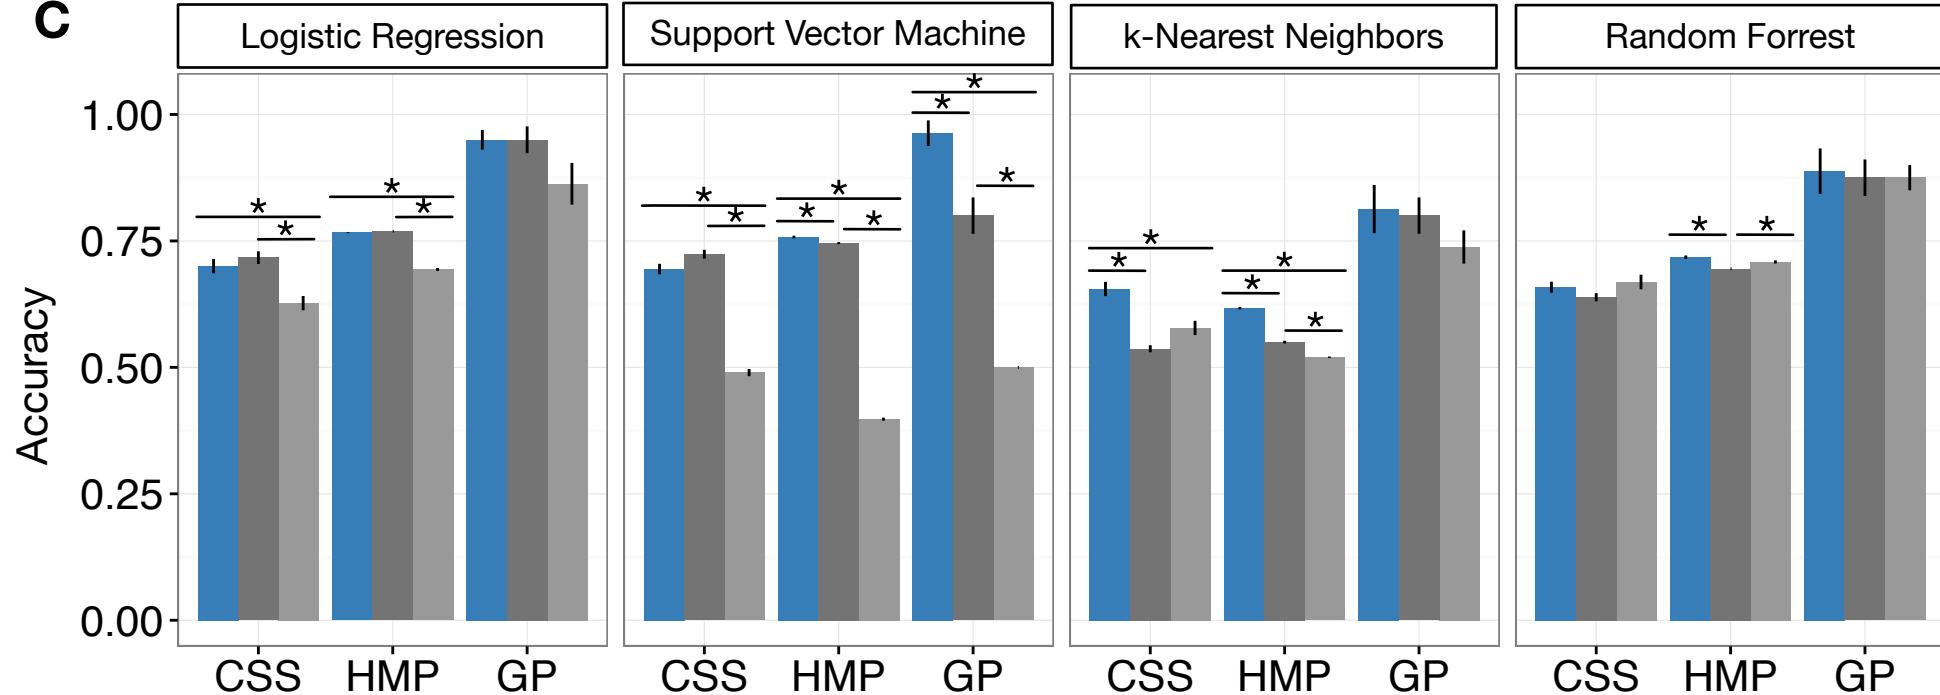
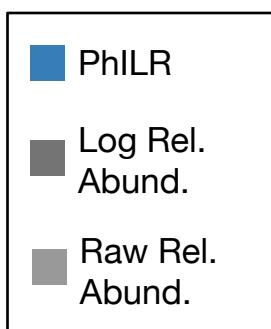
y_1^* Balance of Bacteroides to Ruminococcus and Lactobacillus
 y_2^* Balance of Ruminococcus to Lactobacillus

D Transform in Simplex

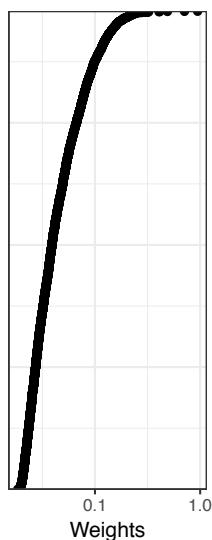
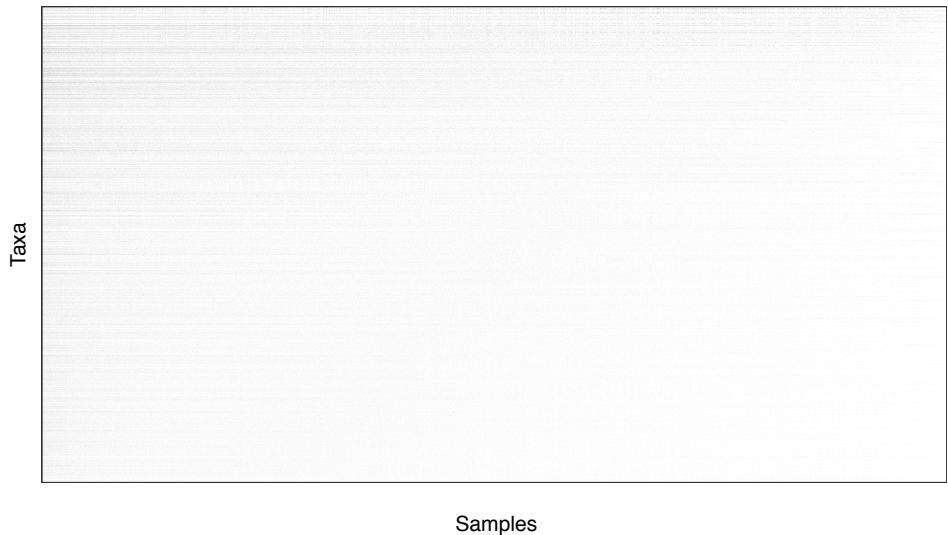
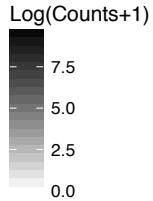


E Data Embedded in PhILR Space

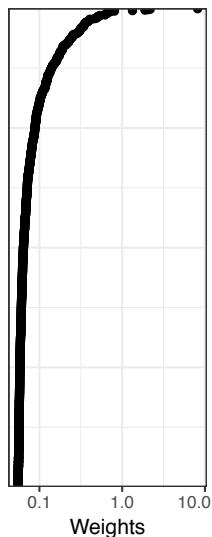
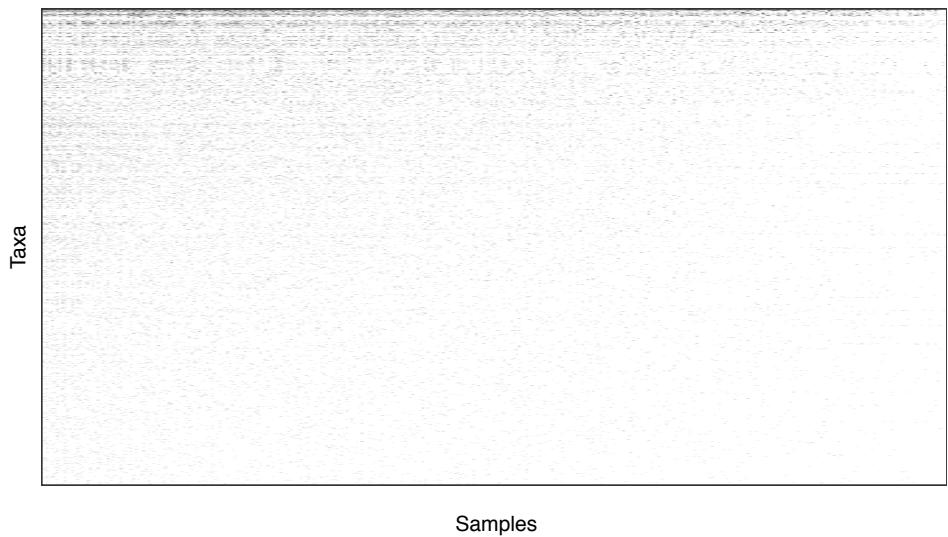
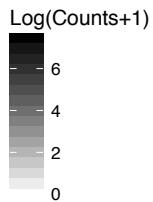


A**B****C**

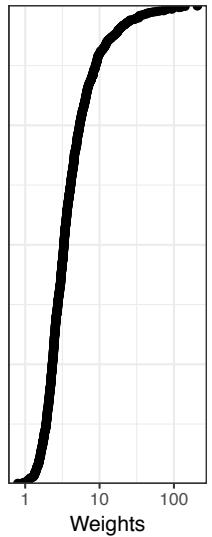
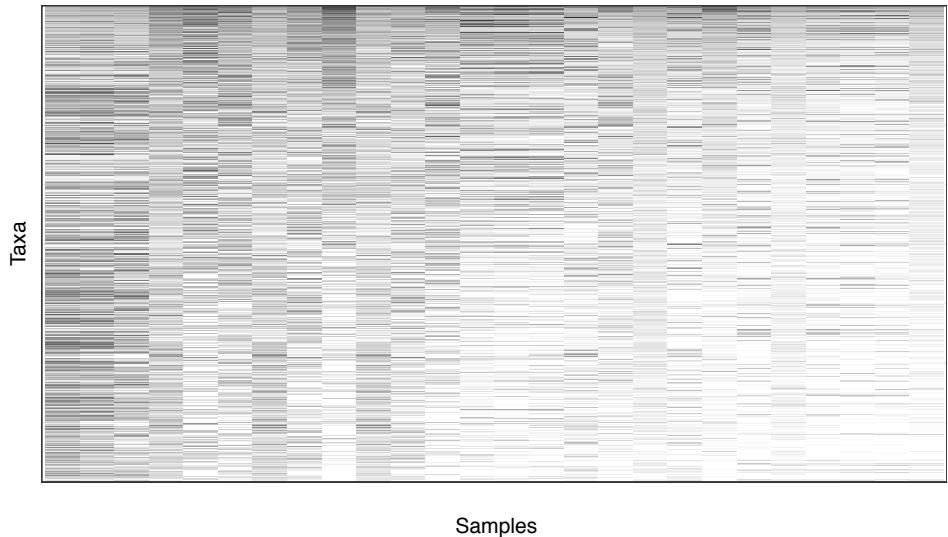
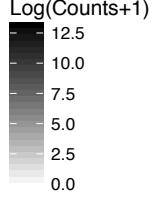
Human Microbiome Project

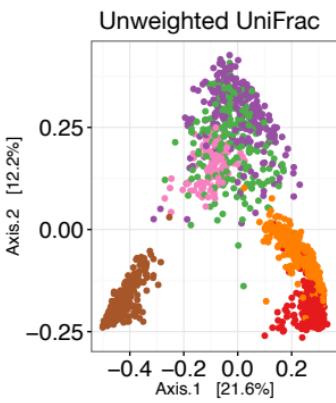
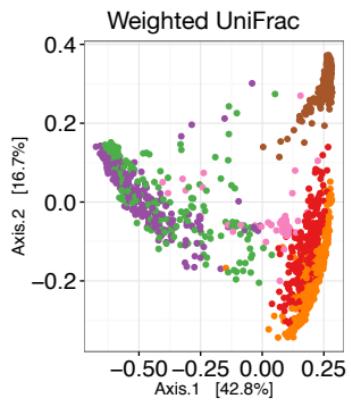
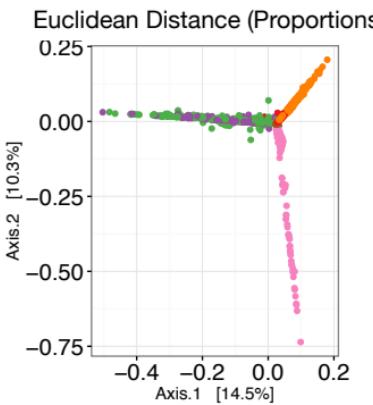
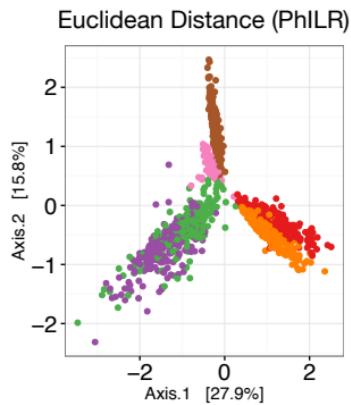


Costello Skin Sites



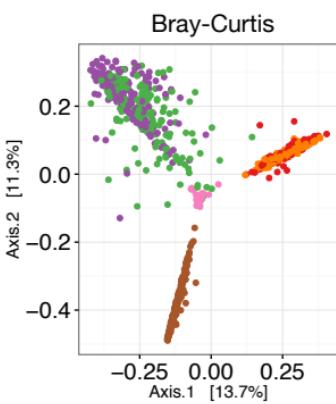
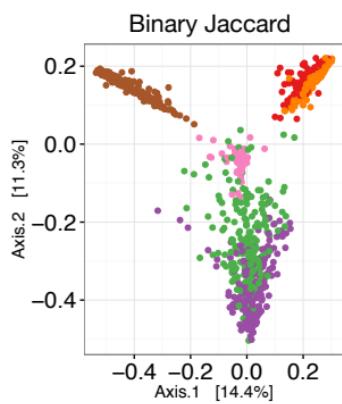
Global Patterns

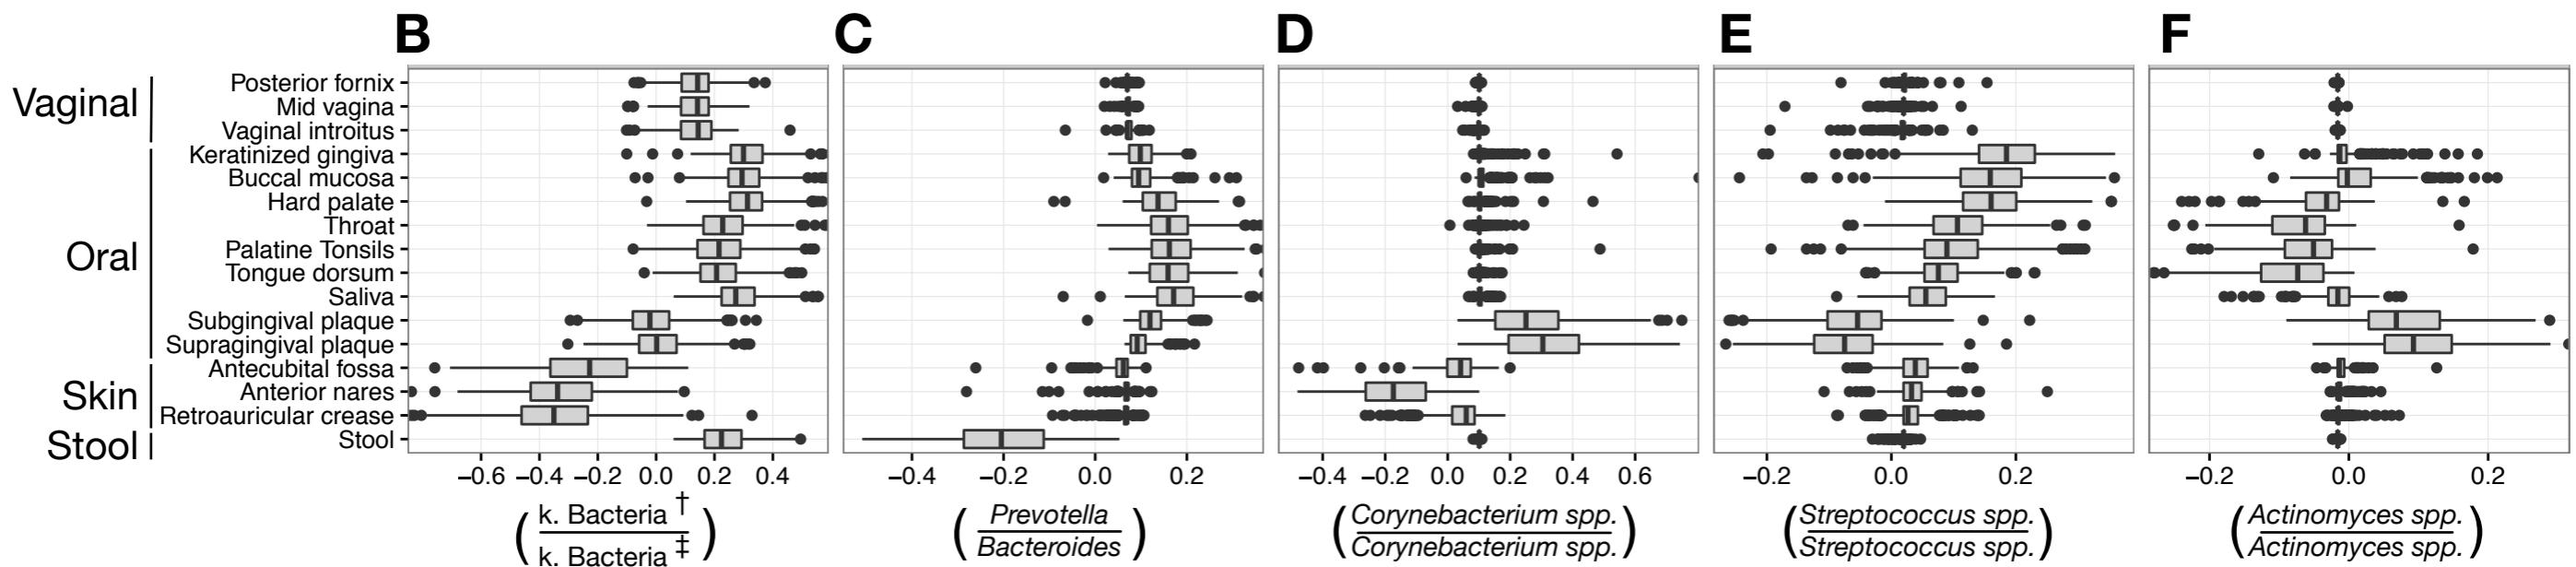
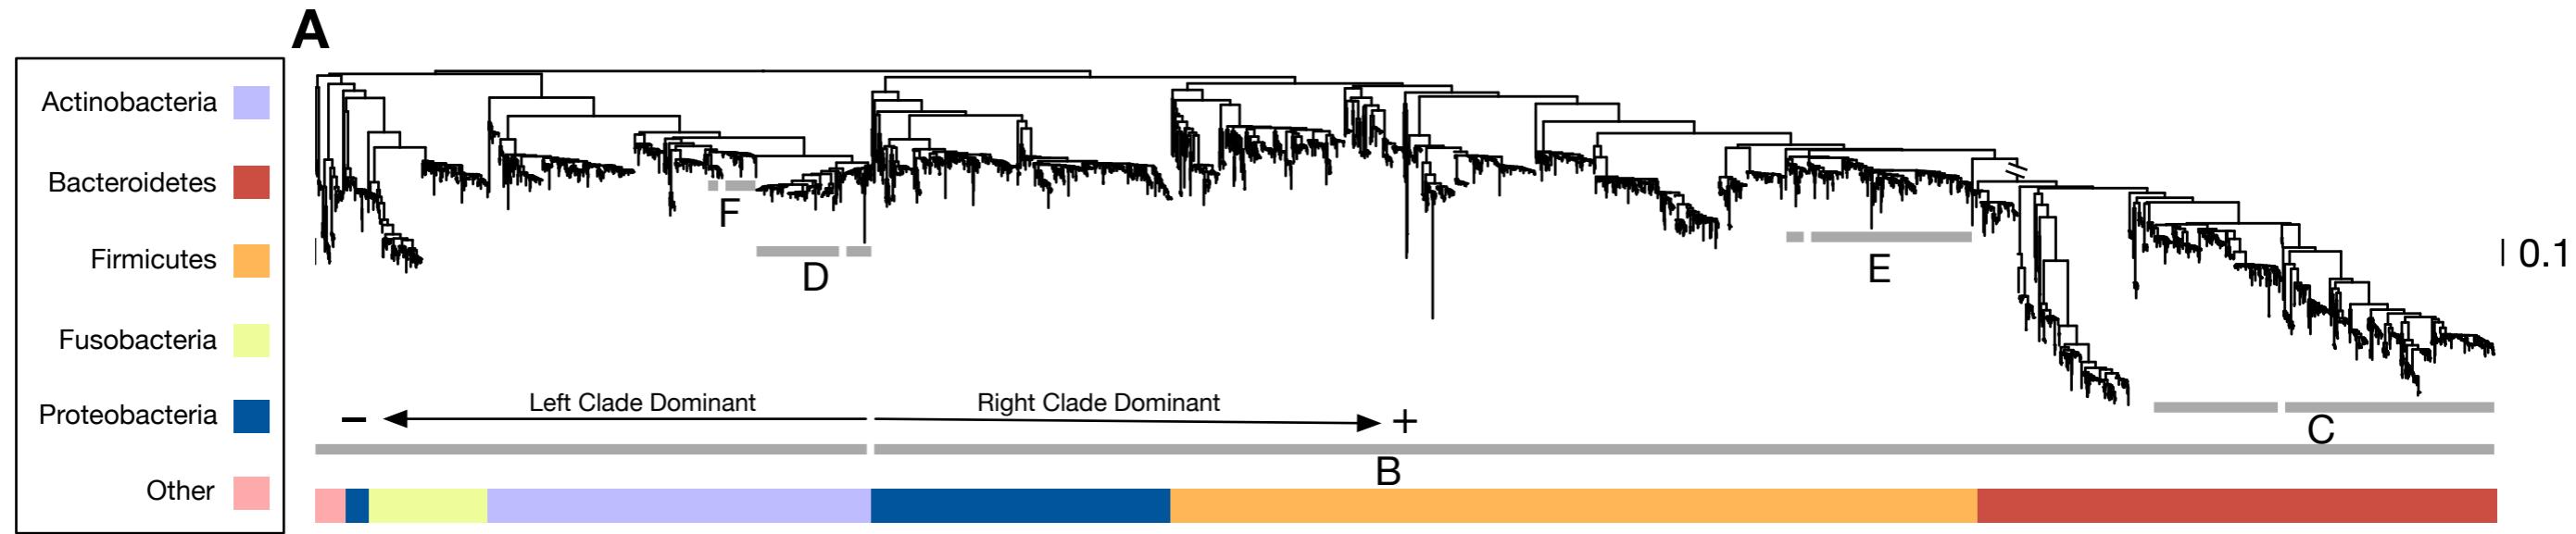


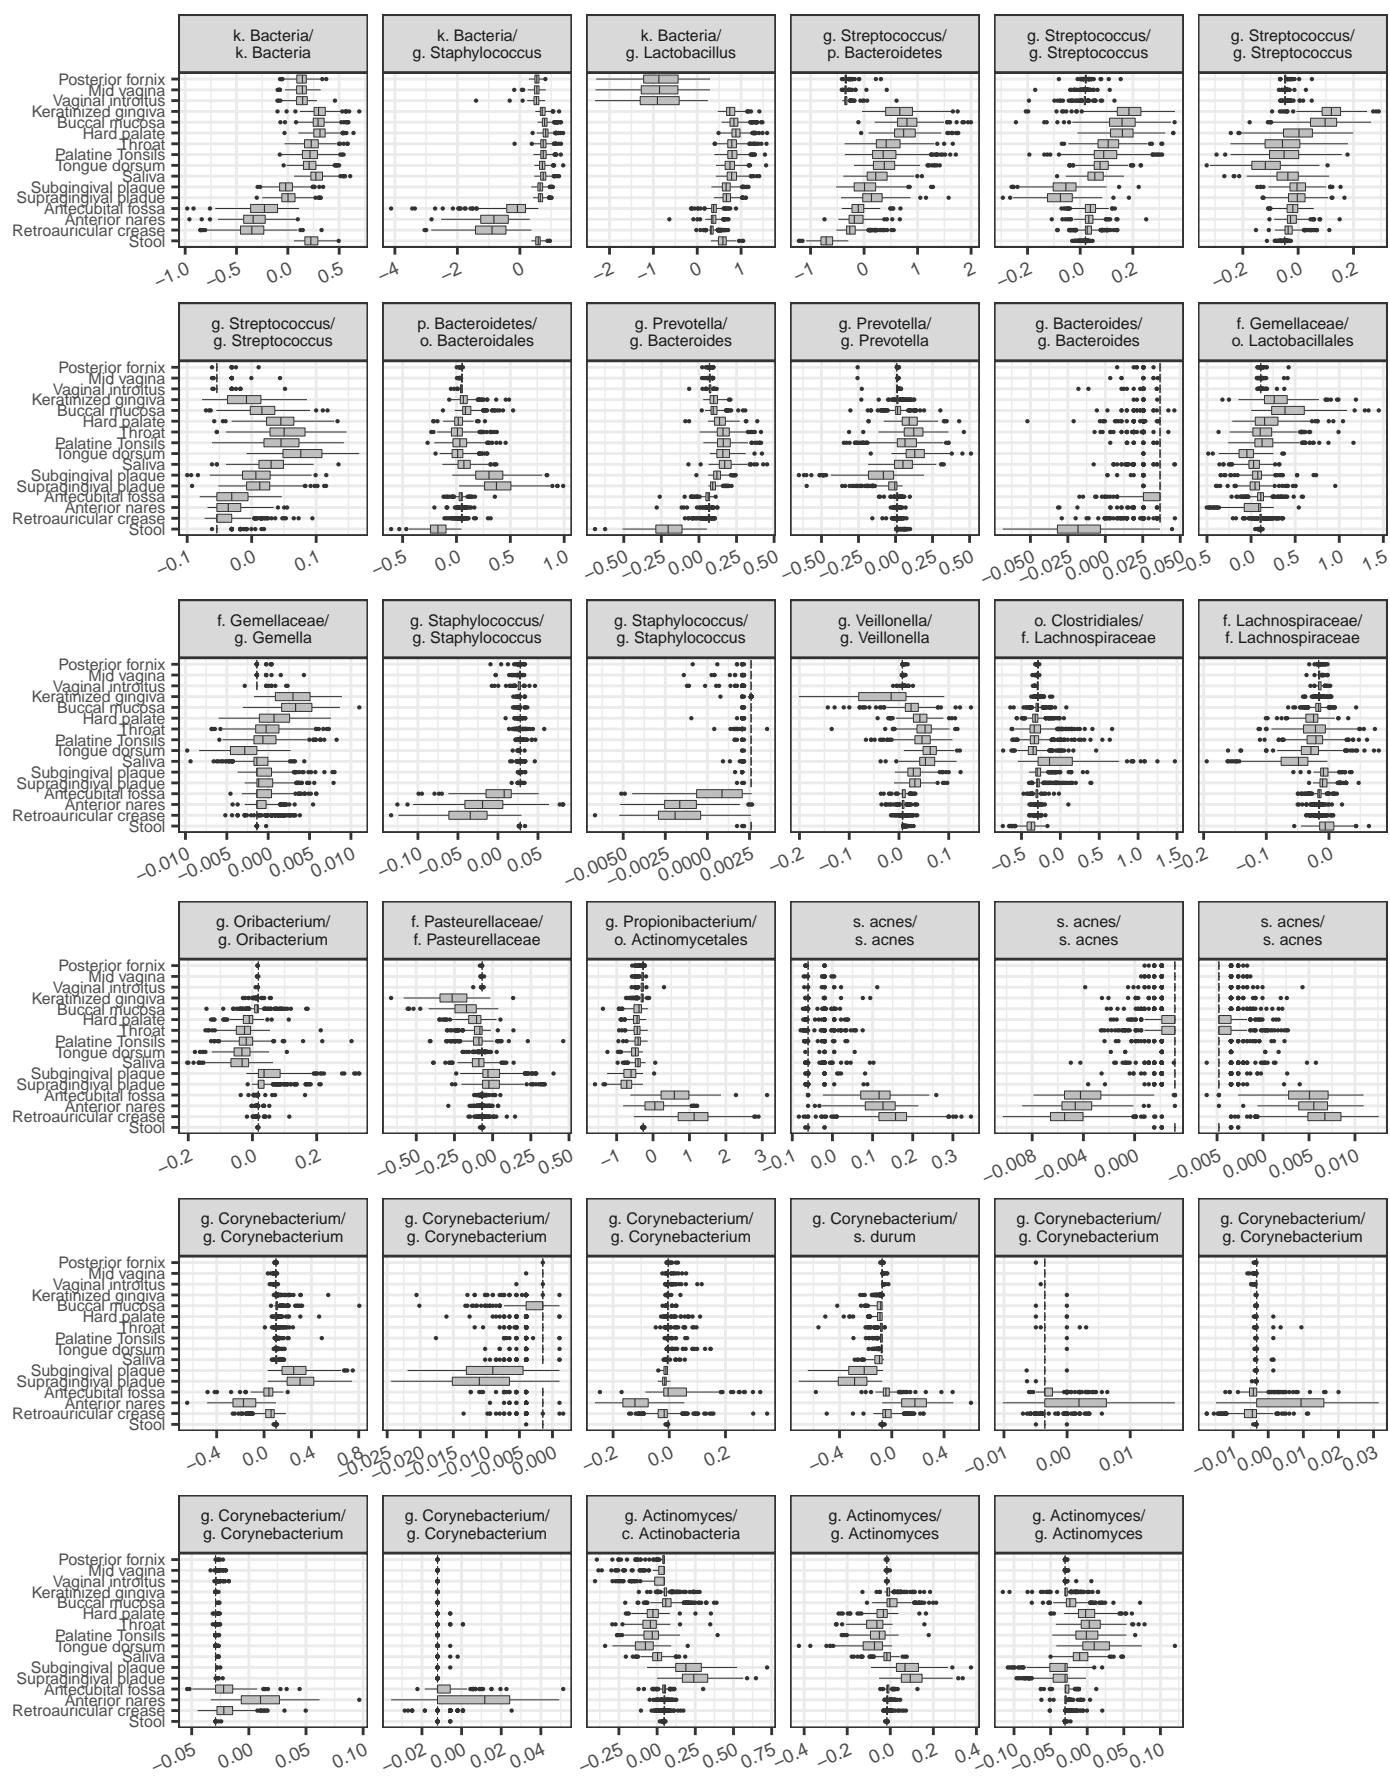


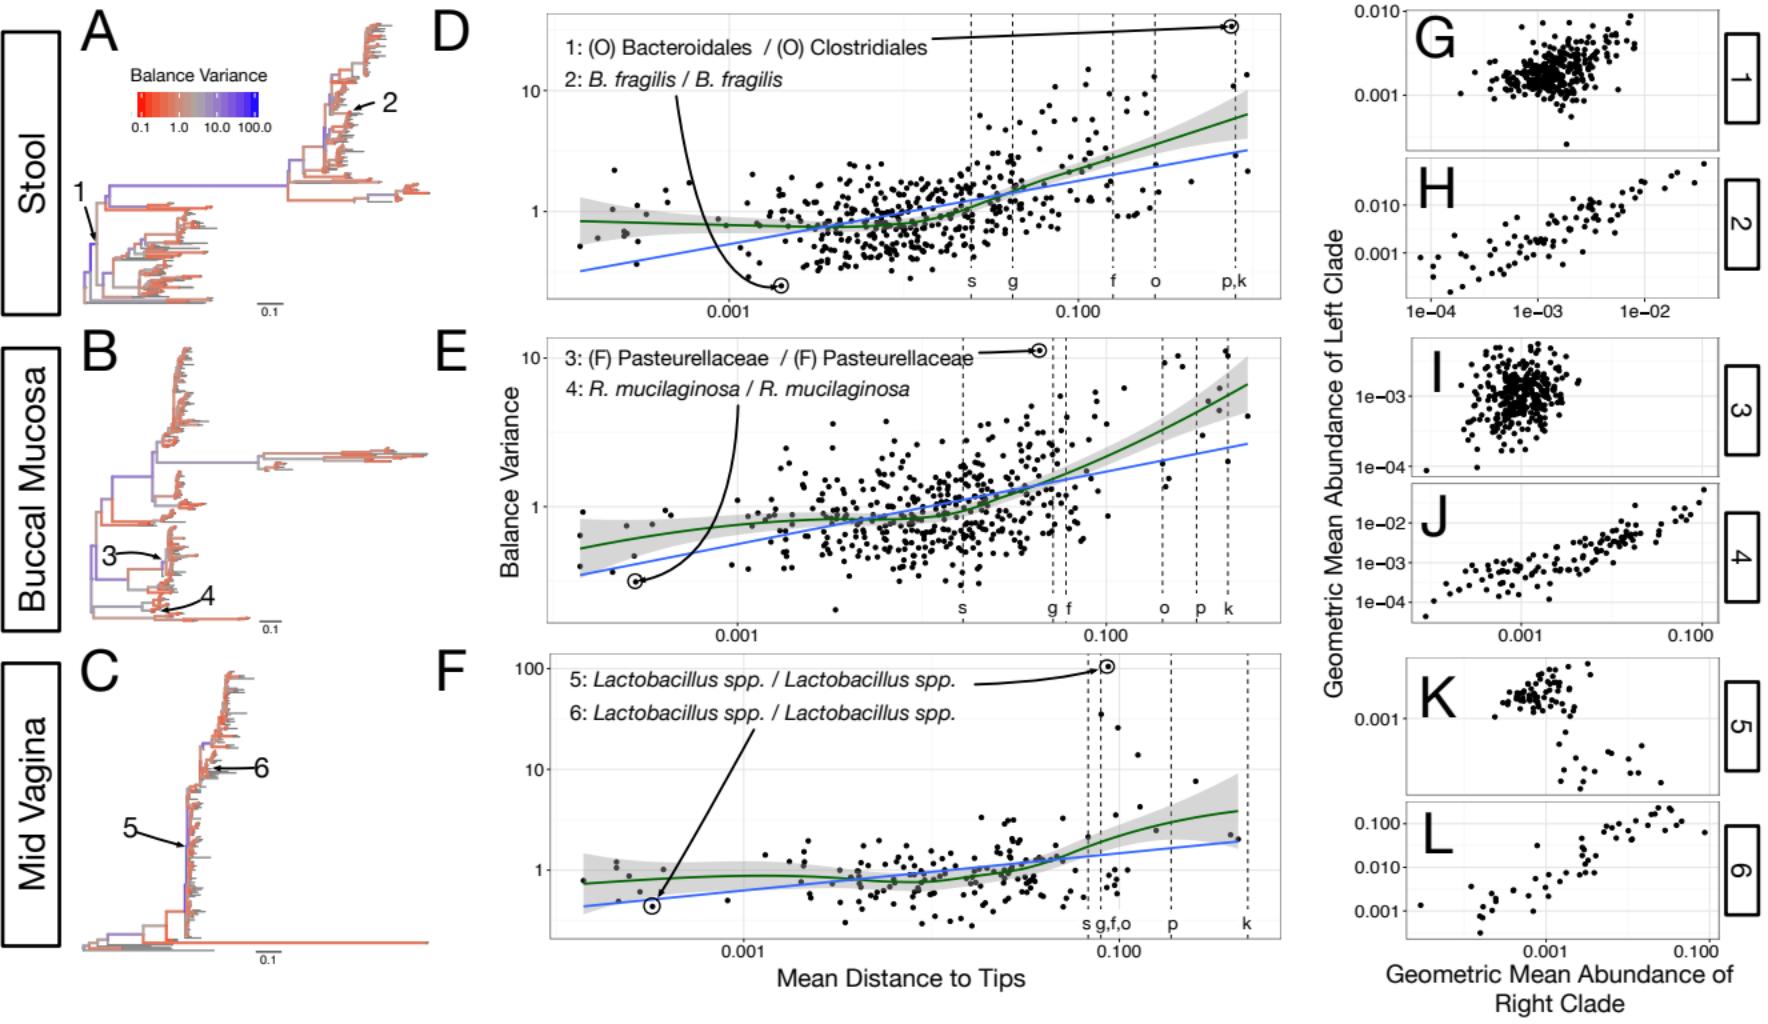
Legend:

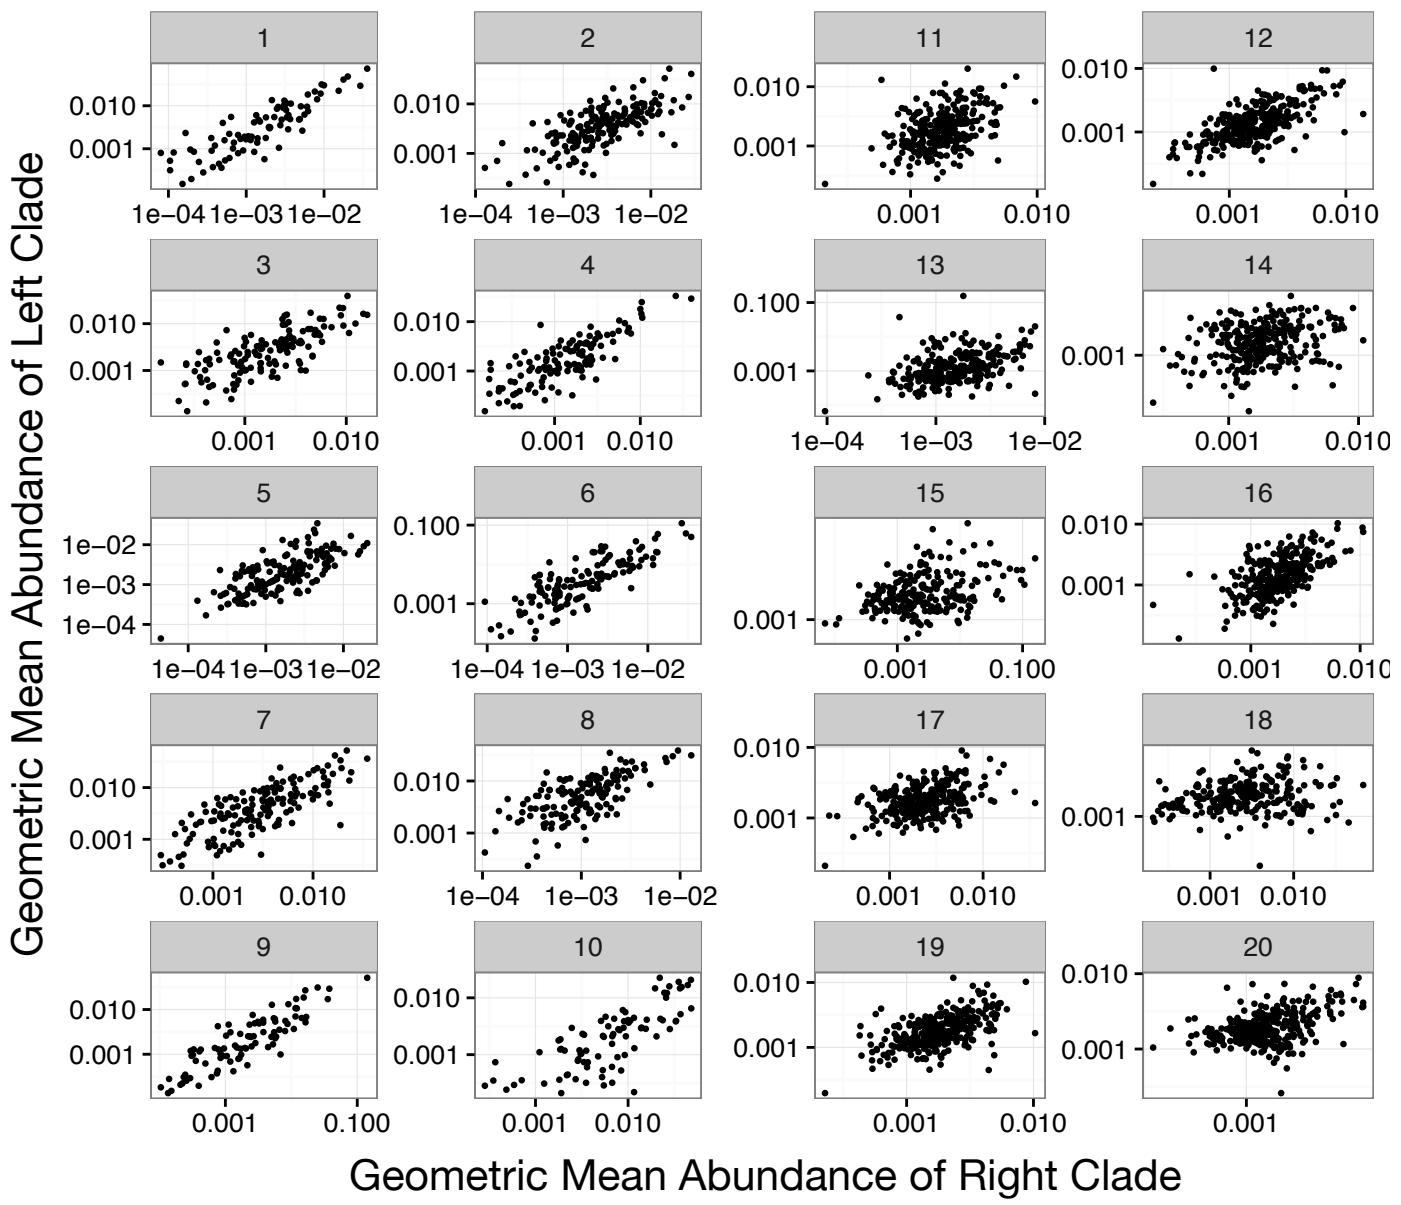
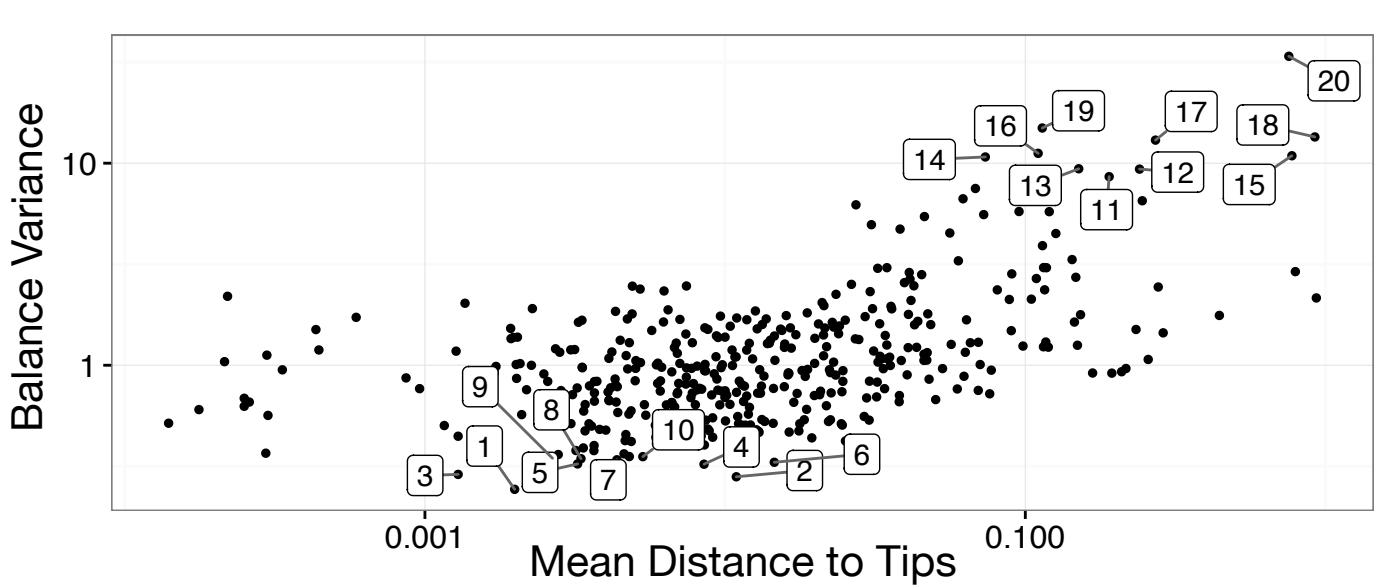
- Antecubital Fossa
- Anterior Nares
- Keratinized Gingiva
- Saliva
- Stool
- Mid Vagina

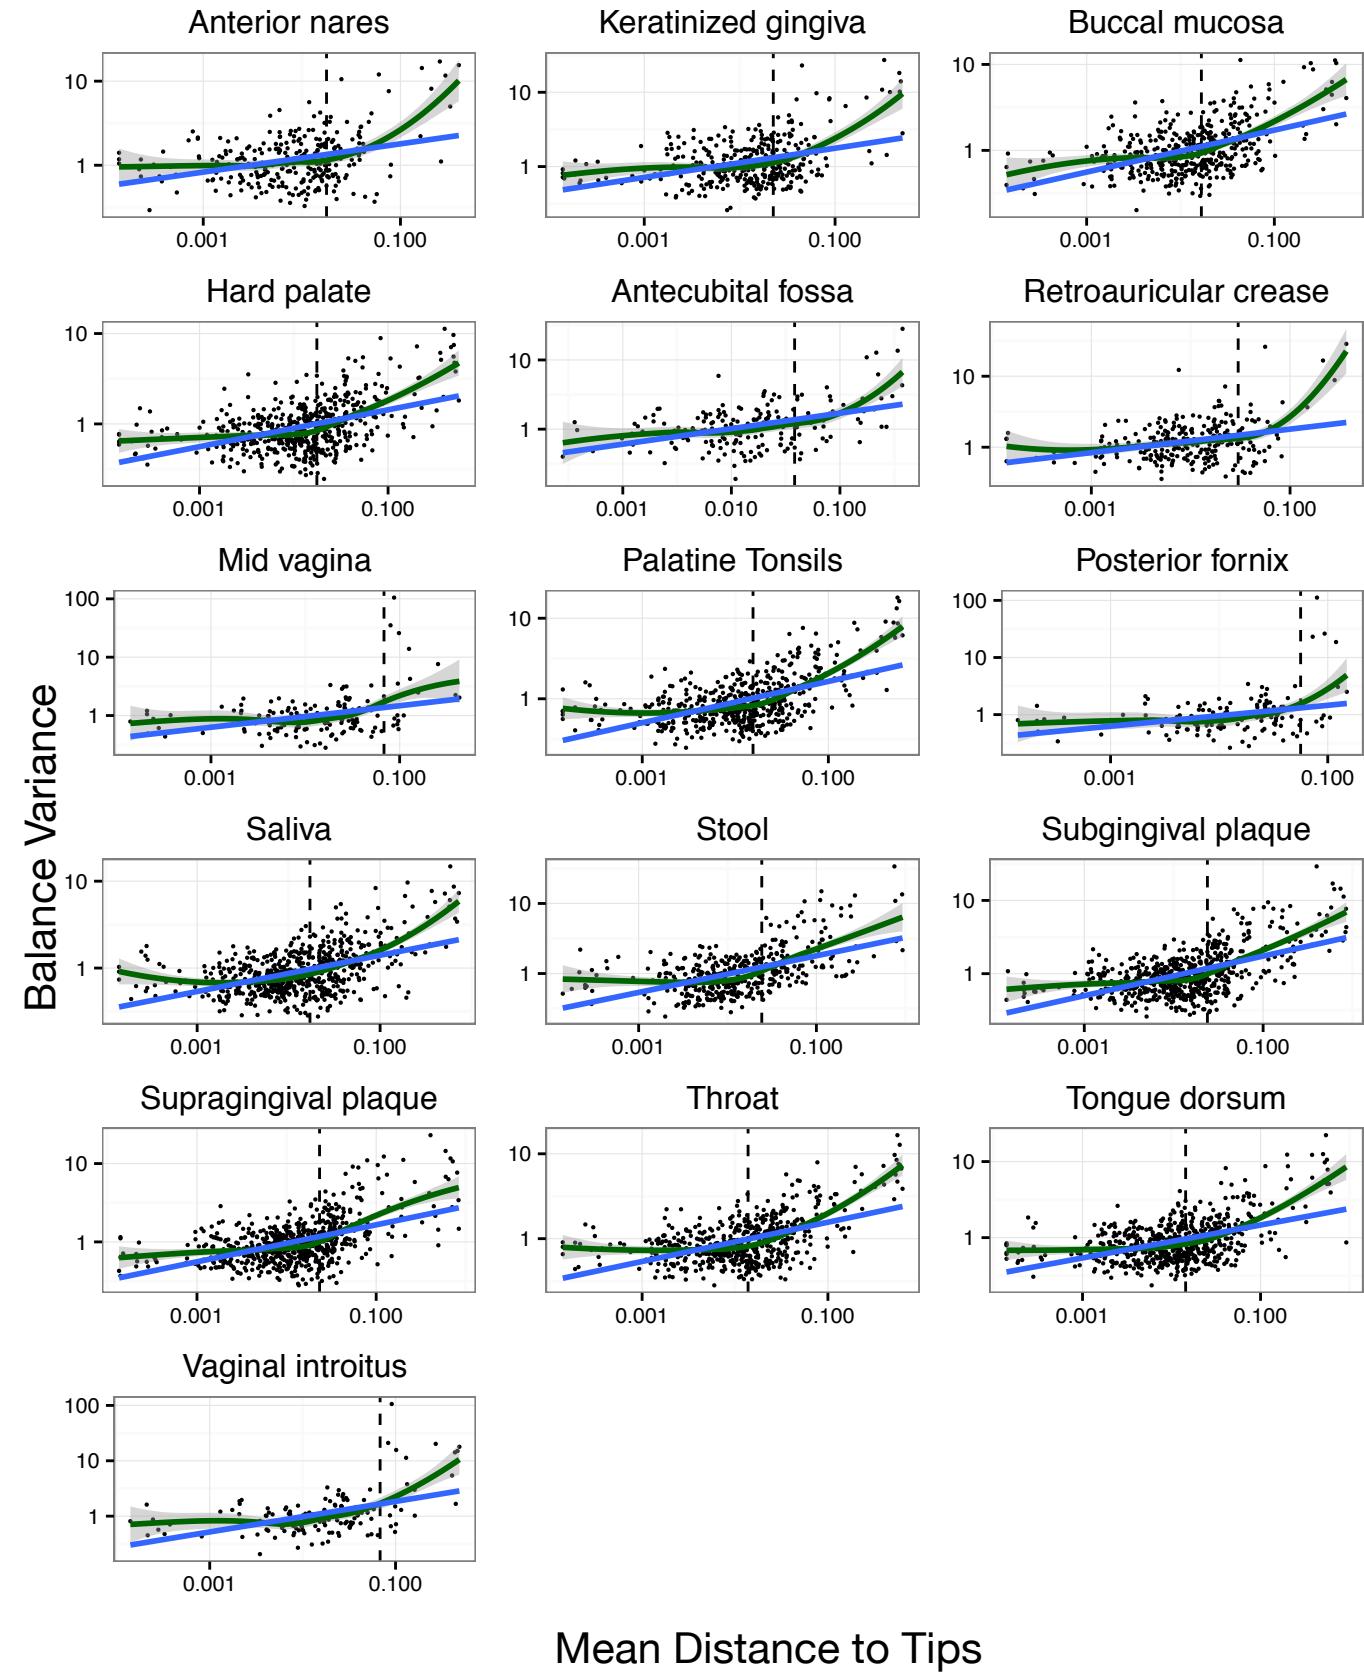




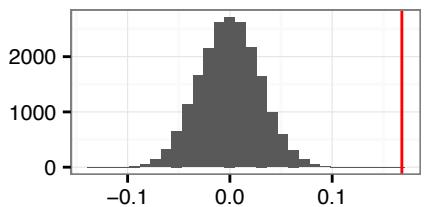




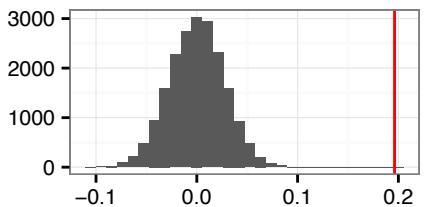




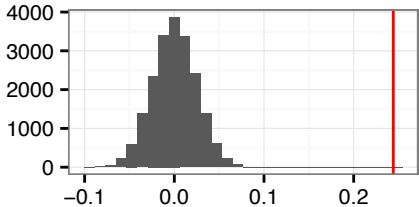
Anterior nares



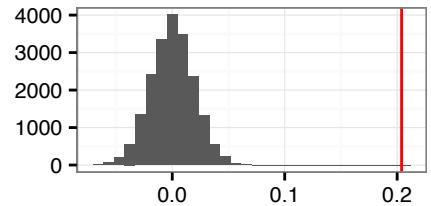
Keratinized gingiva



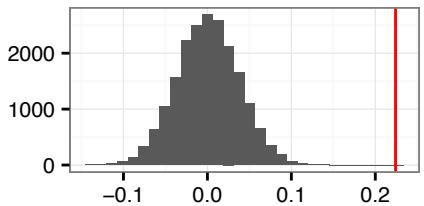
Buccal mucosa



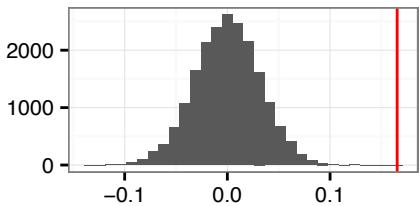
Hard palate



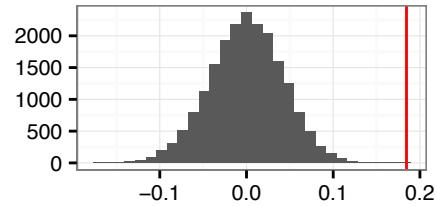
Antecubital fossa



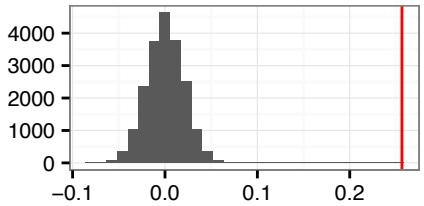
Retroauricular crease



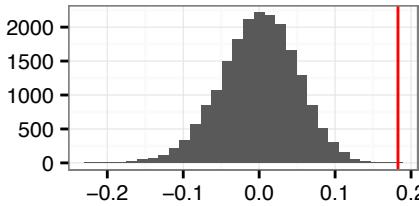
Mid vagina



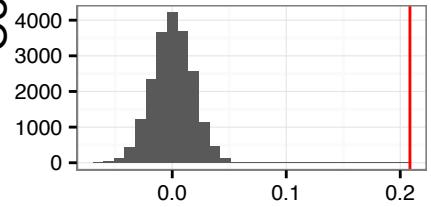
Palatine Tonsils



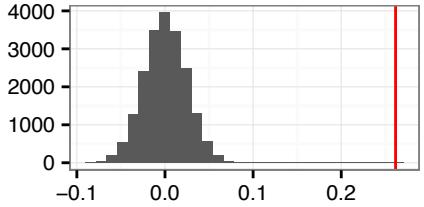
Posterior fornix



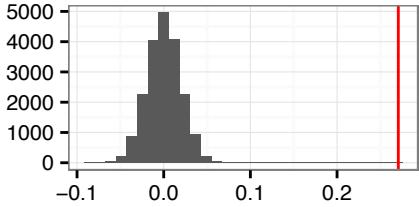
Saliva



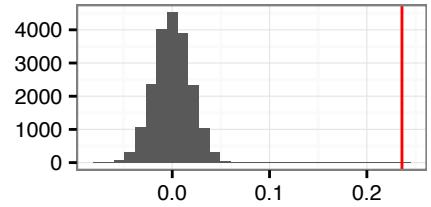
Stool



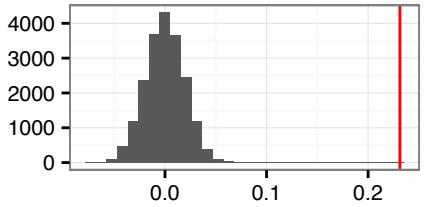
Subgingival plaque



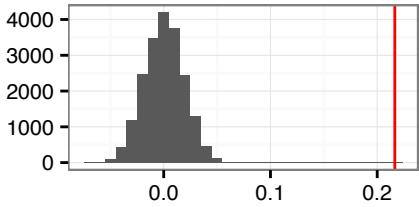
Supragingival plaque



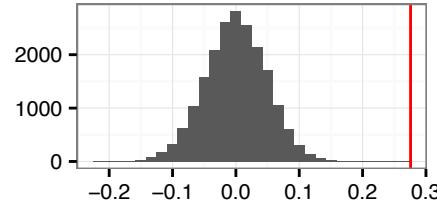
Throat

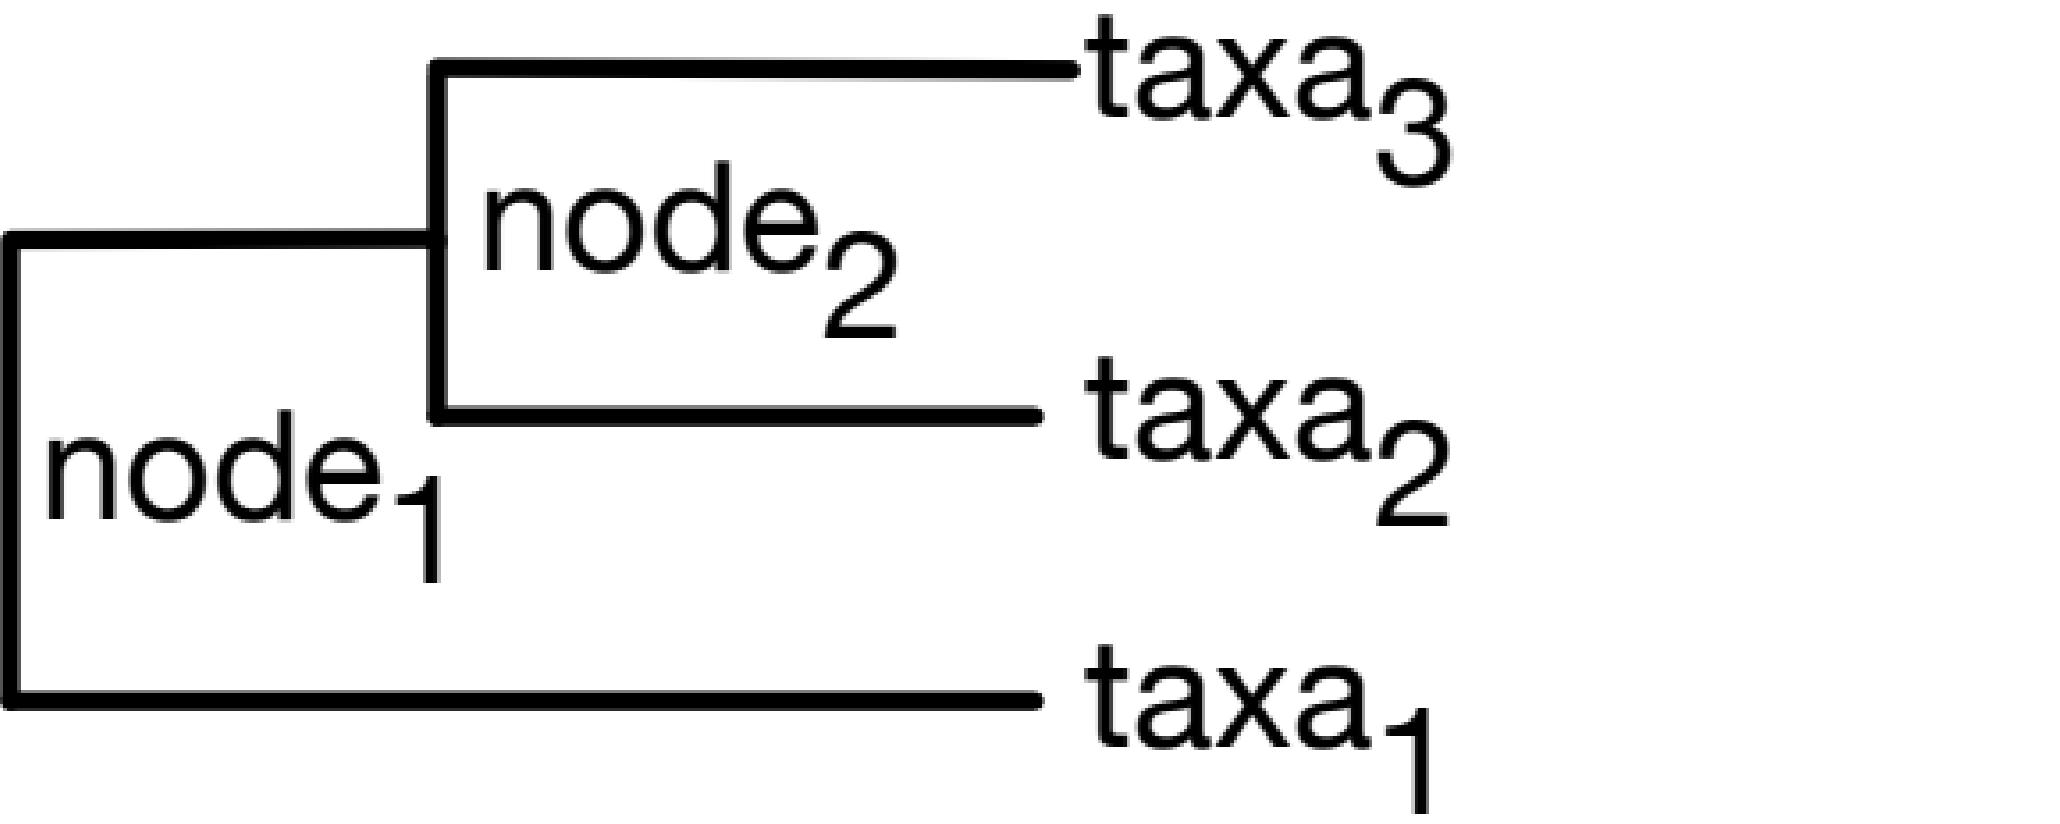


Tongue dorsum



Vaginal introitus

 β



	taxa ₁	taxa ₂	taxa ₃
node ₁	+1	-1	-1
node ₂	0	+1	-1